

Vrije Universiteit Amsterdam



Bachelor Thesis

Implementation of the improved Peterson Kearns rollback algorithm

Author: Robbert Alexander Rutte (2655858)

1st supervisor: W.J.Fokkink
2nd reader: supervisor name

*A thesis submitted in fulfillment of the requirements for
the VU Bachelor of Science degree in Computer Science*

May 16, 2024

Abstract

In a distributed computer network it is important that the system is resilient and that the data integrity is maintained, furthermore it should be able to handle concurrent failures and have a consistent state when ending execution. One of the protocols that satisfies these criteria is the Peterson-Kearns Rollback Recovery protocol. The algorithm is based on a vector-based clock to be able to establish a causal order of the messages.

In this thesis the improved Peterson-Kearns Rollback Recovery protocol [1] will be implemented and experimented with. It includes a background study, discusses the limitations and challenges for the implementation of the protocol, implements the protocol and measures its overhead and efficiency. The experimentation will include both the memory overhead, communications overhead and the run-time complexity of the implementation.

1 Introduction

In the field of distributed computing, a complex system of processes that participate to accomplish a shared goal, inter-process communication (IPC) is indispensable asset. Yet as crucial as inter-process communication is, it also brings a surplus of complications and complexities that require thorough consideration and planning. From data consistency problems, ensure system resilience to minimizing communications overhead.

The protocol discussed in this thesis aims to tackle the data consistency and system resilience problems. The improved Peterson-Kearns Rollback recovery protocol is able to handle multiple concurrent failures without the program ending in an inconsistent state. In the next section the protocol will be described.

1.1 Improved Peterson Kearns rollback algorithm

The protocol provides data integrity and system resilience through the means of of periodically checkpointing the IPC messages to stable storage. however even when operating in a FIFO channel environment a failure of a process could result in an inconsistent state. For a message's send event could be lost to a failure if it is send causally after the last checkpoint.

In the algorithm optimistic recovery is considered, witch indicates that the checkpointing of the nodes is not synchronized. the rollback procedure handles the uncoordinated nodes reach a consistent state. It will do so by handling the *lost messages*, messages where the receive event was lost in a failure at a crashed process or was rolled back to achieve a consistent state, and *orphaned messages*, messages where the send event was lost in a failure at a crashed node or was rolled back. Once the process initiates the restart or rollback procedure the execution is postponed and the messages that arrived are buffered to be handled once the process has restarted or rolled-back.

The protocol relies on a vector based clock that is used to determine the causal order of the messages. To ensure the consistency of the system the Peterson-Kearns protocol relies on a logical vector clock to calculate witch events are causally after events that are lost in a failure. each of the processes will maintain a clock vector and a failure vector to ensure the the ability to uniquely identify failures. The failure vector ensures a consistent state when processes fail concurrently. An in-depth description of the protocol will be provided in Section II

1.2 Context

Inter-process Communication is a vital part of computer science and facilitates an operating system [2] or a distributed system [3] a way to let different processes communicate and cooperate to accomplish a task. While IPC is inescapable to accommodate cooperation and communication between processes it also numerous problems such as reliability and fault tolerance. A system must be resilient enough to deal with failures to ensure consistency and reliability. The performance overhead must also be optimized since processes that use IPC incur a overhead in computational and memory resources. Furthermore IPC is also inherently more vulnerable to security threats [4], this must be managed to ensure the confidentiality and integrity of the data. A system must also be synchronized to avoid race conditions or a system must have ways to deal with race conditions. If a system uses shared memory across processes this must be well managed to ensure its correctness. IPC is a well established in computer science.

1.3 Problem statement

Distributed systems can always be affected by failures, like processes crashing, software errors or outside factors. All of these will be affecting the system consistency and correctness of the algorithm being executed. To prevent the failures from impacting the system implementing a rollback procedure is crucial. The rollback procedure will ensure that any event that is affected by a failure will be undone and forces the system to resume execution from a previously checkpointed consistent state. This will ensure the correctness and the consistency of the system.

1.4 Objectives

The objective of this thesis is to implement the improved Peterson-Kearns rollback protocol. The protocol should facilitate the system to:

1. Enable its processes to notify others of a failure that occurred.
2. Be able to respond to these failures.
3. Rollback the events affected by the failure.
4. Resume execution from a consistent state.
5. Provide an efficient way to implement the procedure
6. Minimize the communication and performance overhead of normal execution.

1.4.1 Scope

The scope of this thesis includes:

- Implementing the improved Peterson-Kearns protocol.
- Testing the procedure under various scenarios to ensure the correctness of the implementation.
- Assess the performance, memory overhead and the communication overhead.

This Thesis will provide:

- The implementation of the procedure with the source code and documentation.
- The test cases used to evaluate the correctness of the implementation.
- The measurement results and evaluation of the results.

- Indicate areas that could be optimized and enhanced based upon the measurements.

Constraints

- The rollback procedure is designed to minimize the performance overhead, since this is an optimistic recovery algorithm, under normal operation of the implementation.
- Minimize the memory overhead of the implementation.
- Mimic the *read()* and *write()* c functions to facilitate easy usage of the implementation.
- The rollback procedure should maintain a consistent state.

In section 2 a study of the background of rollback algorithms and related work will be done. Afterwards in section 3 the improved Peterson-Kearns protocol will be explained in detail, the design choices will be explained, the limitations of the implementation will be specified and the challenges of the implementation will be mentioned. In the fourth section the methodology will be described, the research goal and the implementation details will be provided. The results of the experimental findings and analysis of these results will be done in the 5th section as well as an indication where more improvements could be made. The conclusion will be located at the 6th section.

2 Background

Rollback algorithms originate from the database field of computer science, they were first employed to ensure the consistency of a database system and to assure the atomicity and persistence of transactions when failures occur. The types of rollback algorithms that exist are categorized in two segments: *Immediate rollback*, where the algorithm rolls back the affected events immediately after detecting that a failure has occurred and *Deferred rollback*, where the algorithm defers the rollback of affected events and handles them at a later time.

Checkpointing is the periodically saving of events to stable storage. Checkpointing is a central component in every rollback algorithm since when a failure occurs that makes a process crash, it must be reinitialize with part of the memory

otherwise to maintain a consistent state the algorithm would need to rollback all events. Therefore to be able to resume from some earlier consistent state checkpoints of the events that occur at all processes need to be saved to stable storage. Checkpoints frequently include the state of the system and previously generated events.

Logging is another critical component of rollback algorithms. A log is kept of the send and receive events that are processed by the process. This log is used when a failure occurs at another process to undo the changes that were made by events that are affected by the failure and restore the state to a consistent state. Logging frequently includes the execution order of the events.

Coordinated checkpoints is where all processes of a system coordinate the checkpointing. In that case the algorithm synchronizes the checkpointing of its processes. They are often used of a way to maintain a consistent global state. This enables the protocol to rollback to a certain checkpoint with all its processes which will achieve a consistent state. No further processing of the checkpointed messages will be necessary since they will not exist.

Since the inception of rollback algorithms they have been applied in many fields other than that of database systems, like in distributed systems and operating systems. Rollback algorithms are vital in databases, operating systems, distributed systems and other fields to handle failures and preserve a consistent state.

2.1 Related Work

in this subsection we will explore existing rollback algorithms and classify them as immediate rollback or deferred rollback.

Immediate rollback

An immediate rollback algorithm is a protocol that will apply changes to the system immediately in addition it will maintain a log file where the applied changes are stored. This often involves committing procedures for every event to ensure the consistency of the system. Once a failure occurs the rollback procedure will undo the changes to the system that have not been committed yet. The major advantage of this method is that the recovery process will not be an expensive procedure, thus this method is usually preferred when failures occur often in a system. Examples of the immediate rollback algorithms are:

Immediate update described in Principles of Transaction Processing [5]. Where an event is either committed when it executes or it is rolled back immediately. These commits are coordinated with other processes. Once a failure occurs the current event is rolled back to ensure a consistent state.

Compensation-based recovery as described in Sagas [6]. In this algorithm when an events affects the state of a system a compensation event is stored in memory. Once a failure occurs the compensation event is executed to revert the state of the system to a consistent state.

Deferred rollback

A deferred rollback algorithm is a protocol that delays the commit procedure of events. This involves the storing of events in a log. Once the commit procedure is complete the changes to the system stored in the log will be applied. Once a failure occurs the rollback procedure will remove affected events in the log and these events will not be affecting the system. The advantages of this method is that the I/O messages will be kept low since multiple events are committed at once, thus this method is usually used if failures will only occur periodically. Examples of deferred rollback algorithms are: *Vector based rollback-recovery* in Rollback based on vector time [7]. This is the original algorithm that the improved Peterson-Kearns recovery-rollback [1] protocol is based upon. In this algorithm causal order of events is provided by a logical vector clock. Once a failure occurs other processes are notified via tokens and the affected events are rolled back to maintain a consistent state.

Distributed transaction processing as described in Checkpointing and Rollback-Recovery for Distributed Systems [8]. In this algorithm a two-phase-commit procedure is used to facilitate the rollback. In the two-phase-commit procedure all processes coordinate the checkpointing to ensure a consistent state.


3 Framework

In this section an in-depth overview of the implemented protocol will be given. This will be a summary of the information given in the original Msc thesis of van Eck,C [1].


3.1 Classifying events

In this subsection all events will be classified. An event is either a sent or a received message both classified as a send event and a receive event respectively. When these events are affected by a failure they will become either an orphan event or a childless event, these will be further classified into orphan-arrived, orphan-in-transit, childless-lost and childless-orphan events.

3.1.1 Orphan event



(a) Orphan-arrived




(b) Orphan-in-transit

Figure 1: Orphan events


An *orphan event* is created when the send event of a receive event is lost in a failure, this lost send event will be causally after the last checkpoint at the failed process.

- **Orphan-arrived:** The receive event is causally before the rollback is performed, and is thus before the control message signaling a failure has arrived. Event r_1^1 in figure 1a becomes an orphan-arrived event after the failure at p_0 .
- **Orphan-in-transit:** The receive event is causally after the rollback has been performed, and is thus after the control message signaling a failure has arrived. Event r_1^1 in figure 1b is a orphan-in-transit event.

3.1.2 Childless event



(a) Childless-lost



(b) Childless-orphan

Figure 2: Orphan events

A *childless event* is created when the receive event of a send event is lost in a failure, this lost receive event will either be itself lost in the failure or will be lost in the rollback of another process.

- **Childless-lost:** The receive event that is lost is causally after the last checkpoint at the failed process, and is thus itself lost in the failure. Event s_1^1 in figure 2a becomes a childless-lost when the failure at p_0 occurs.
- **Childless-orphan:** The receive event that is lost is causally after an orphaned event at another process, and is thus lost in the rollback procedure at the other process. Event s_2^1 in figure 2b becomes a childless-orphan after the failure at p_0 occurs.

3.2 Algorithm Overview

The improved Peterson-Kearns Rollback recovery protocol [1] has several procedures that are needed for the protocol to function correctly. The protocol employs vector clocks to determine what events are affected by a failure. Each process keeps track of 2 vector clocks, a clock for keeping track of both casual order of the events, the time

vector, and the number of times a process has crashed, the failure vector. Both of these clocks will be sent along with a message to inform other processes of the current state of the vector-clock at the sending process. All events will be stored in volatile memory along with its accompanying vector clocks.

3.2.1 Time Vectors

Each process has an associated time vector. The time vector $T = T^0, \dots, T^{n-1}$ gets updated whenever an event e_i occurs where n is the number of processes and i is the process id where the event occurred. The T' notation is the time vector before the event e_i occurred.

- When e_i is a send event then the time vector T of a process gets updated to $T^i = T'^i + 1$.
- When e_i is a receive event the time vector T of a process gets updated to $T^j = \max(T'^j, e_i.T^j)$ where $j \neq i$ and $T^i = T'^i + 1$

3.2.2 Failure vectors

Each process has an associated failure vector, the failure vector $F = F^0, \dots, F^{n-1}$ gets updated when a failure occurs in a process or the process get notified of a failure at another process.

In the case of the failure failure happening at the current process will perform the restart procedure, and therein the failure vector gets updated to $F^i = F'^i + 1$ where i is the id of the current process, afterwards the process sends control messages to all other processes to notify them of its failure. This control message includes:

- *id*: The process ID of the failed process.
- *n*: The failure number, indicating the amount of failures that occurred in total at the failed process.
- T_i^i : The restart time, indicating the T^i at the moment of restart at the failed process.
- *arrived*: The set of (T^j, F) of the messages received by the failed process that were sent by the process p_j receiving the control message.

In the case that the failure did not occur at the current process, then the current process will be notified of the failure by receiving a control message. When the control message is received the process will perform the rollback procedure,

and therein the failure vector will be updated to $F^j = n$ where n is the value received in the control message for the failure number and j is the process ID of the failed process.

3.2.3 Checkpoint procedure

During a checkpoint procedure all information needed to restore the process to a stable state are stored in stable storage. The information to be saved in storage is as follows:

- *state* : the state of a process.
- T : The time vector of the process at the moment the checkpoint occurs.
- *log*: The log where the events that occurred at the process are stored.

During a checkpoint procedure no other event can occur.

3.2.4 Restart procedure

After a failure at a process it will need to restart from the last checkpoint . The operations needed to perform the restart successfully are:

1. The failure vector is updated to $F^i = F'^i + 1$ where i is the ID of the process where the failure occurred.
2. The process will be restored to the latest checkpoint.
3. Send control messages to the other processes.

During a restart procedure no other events can occur.

3.2.5 Rollback Procedure

After a control message is received at a process it will need to perform a rollback to achieve a global consistent state. The operations needed to perform the rollback successfully are:

1. The failure vector is updated to $F^i = F'^i + 1$ where i is the ID of the process where the failure occurred.
2. If some event is orphaned at the process that received the control message, then the state is returned to the last checkpoint that does not include orphaned events. All events after this checkpoint that are not orphaned will be replayed.

3. Any receive events that have been orphaned but have a corresponding send event that was not orphaned will have their receive event replayed.
4. All messages that this process sent to the process where the failure occurred with a non orphaned send event that did not arrive will be re-transmitted.
5. A new checkpoint will be created.

During the rollback procedure no other events can occur.

3.3 Limitations

The limitations of this algorithm are that the maximum number of messages that can be sent by one process will be the upper-limit of the integers used in the time vector, this could be mitigated by allowing the time vector values to roll over back to zero in some way but this would require more logic to accommodate for this.

Another limit of this algorithm would be that when the amount of events stored by a process becomes large and a failure occurs at that process it will need to send a large set of messages that have arrived at that process. This could be mitigated by occasionally checking what messages have been checkpointed by all processes and then committing these events. After the set of arrived messages will not be needing messages in it that have been committed. This will be explained more in-depth.

If a task needs to be completed that will generate a large amount of event the volatile memory of a process could reach the allowed limit fast. to mitigate this some events could be removed if they have been committed and will not be needed for re-transmission. This would shrink the volatile memory needed for continuous operation. This will be explained more in-depth.

3.3.1 Commit procedure

When a commit procedure is initiated by a process commit messages will be sent between processes. A committed event will be stored in a separate file in stable storage. These messages will be as follows:

1. the process initiating the commit procedure will signal all other processes.

2. When this signal is received at a process it will send the initiating process the T^i of its last *checkpoint* where i is the process ID.
3. After all of the aforementioned messages have been received the initiating process will construct a minimal checkpointed time vector from these values with its T^i of its last *checkpoint*.
4. All events that have a time vector that is \leq the minimal checkpointed time vector can be committed.
5. The initiating process will send this minimal checkpointed time vector to the other processes. Along with the list of committed receive events that have been sent by the recipient of this message, this is needed for the removal of send events.
6. All processes will now be able to commit events, if the process is not the initiator of the commit procedure it will send all other processes a list of committed receive events that have been sent by the recipient of the message, this is needed for the removal of send events.

in between commit messages other event can occur.

3.3.2 Remove procedure

After all of the other processes have sent the list of committed receive events to a process it can perform the remove procedure. $NextCk(e)$ is the checkpoint after the one event e is stored in. the events that can be removed are as follows:

- *Messages stored in volatile storage*: all events e that have $NextCk(e) \leq$ the minimal checkpointing time vector can be removed. if e is a sent event its corresponding receive event needs to be contained in the set of committed messages received during the commit procedure.
- *Checkpoints stored in stable storage*: All checkpoints ck_i can be removed if $NextCk(ck_i) \leq$ the minimal checkpointing time vector.
- *Failures stored in stable storage* : these can only be removed when FIFO channels are used.

3.4 Design Choices

Since one of the requirements of the implementation is that operation without failures occurring should have a small overhead the implementation will be written in C++. In addition to that to ensure fast operation of vector logic the vector package from the C++ standard library was used. To ensure low overhead on the sending and receiving procedures no loops were used in this part of the implementation. The functions to send and receive messages are modeled after the C write and read functions to facilitate simple usage of the implementation.

An implementation will be made for the base protocol without the commit and remove procedures, as well as a version with the commit and remove procedures to test them against each other for a difference in performance.

4 Methodology

4.1 Research Goal

The goal of the experimental setup is that the implementation will be tested on performance. The performance tests that will include the testing performance of both the implementation with and without the commit and remove procedure. For both of these the performance of the sending and receiving of messages will be measured. The performance of the Restart and Rollback procedure will also be measured. The performance of the rollback procedure will also be measured when dealing with concurrent failures from another process. Afterwards the measurements of both of the implementations will be compared.

4.2 Implementation Details

The implementation of the improved Peterson-Kearns rollback recovery protocol is written in C++ version 20 and is provided in appendices A and B.

State class

The State class contains all components needed to maintain a consistent state. the main components of the class are the id of the process, the number of messages in volatile storage, the event log containing the message events, time vector, failure vector, the file descriptors of all other processes (these will be used to resend lost messages),

checkpoints and checkpoint time vectors, a set arrived messages and a set of arrived control messages. The version of the class where committing and removing messages is possible the State class also includes a vector to track if the commit or remove procedures can be executed, a set of committed messages, a set of committed receive events from the other processes and a minimal time vector to calculate the messages that can be committed and removed. The most important methods of the State class are listed below.

Public methods:

- **State:** Initializes the class. If the initialization is executed with the restart boolean set to true it will execute the restart operation and sent control messages to other processes.
- **checkpoint:** Performs a checkpointing procedure .
- **send_msg:** Sends a message and increments the time vector. This will only send messages of type MSG
- **recv_msg:** Receives all message types and handles them accordingly, it will receive normal messages that increments the time vector, it will receive control messages and initiate the rollback procedure, it will receive all commit messages and process them accordingly and it will receive void messages which will not affect the state of the process and will only be passed through to the caller of `recv_msg`. The method uses `c` function `read` to read from its read file descriptor.
- **send_ctrl:** Sends a control message to all other processes to indicate a failure has occurred. The method uses the `c` function `write` to write to the file descriptor of other processes.
- **signal_commit:** Signals to all other processes to start the commit procedure.
- **update_fd:** Updates the file descriptors of a process.

The public methods also include methods to retrieve the stored event log.

In addition to these public methods the class has several private methods that are used by the class to operate. The most important private methods of the State class are listed below:

Private methods:

- **check_duplicate:** Checks if a received message is a duplicate .
- **check_duplicate_ctrl:** Checks if a received control message is duplicate.
- **check_duplicate_commit:** Checks if a message has already been committed.
- **check_orphaned:** Checks if a received message is orphaned.
- **rem_log_entries:** Removes log entries from the event log.
- **rem_checkpoints:** Removes checkpoints.
- **store_msg:** Stores a message in the event log.
- **commit_msgs:** Commits the indicated messages.
- **rollback:** Performs the rollback procedure. This procedure could resend messages that were lost in the failure to the failed process.
- **send_commit:** Sends its own the time vector value to the process that initiated the commit procedure.
- **commit :** Performs the commit procedure and sends the committed receive events to all other processes.
- **remove_data:** Performs the remove procedure.

The private methods also include serialization methods and methods to calculate various pointers and indices. The **send_void** function is included in the namespace `Pet_kea`. It send a message to a process that does not increment the time vector and is not stored in the event log.

the class is able to send the types of messages listed below:

- **MSG:** This is a message that will be processed by the protocol, it increments the time vector of the process. It includes the time vector and failure vectors at the time of sending, the ID of the sending process and the message size in bytes.
- **CTRL:** This is the message that will be send if a failure has occurred, it will increment the failure vector of the process receiving it. It includes the failure log entry, the ID of the sending process, the amount of messages the

sending process has received from the receiving process of the control message and the time vector value and failure vector of these messages. this message can only be sent after a process has restarted.

- **VOID:** This is the message that will not impact either the time vector or the failure vector and will be passed through to the caller of the **recv_msg** method. And thus not affecting the state of the protocol. This would be used when a message has to be sent between processes that can not alter the state.
- **COMM1:** This is the message that will be sent to initiate the commit procedure
- **COMM2:** This message will contain the its time vector value of its latest checkpoint.
- **COMM3:** This message will contain the combined time vector values of the latest checkpoints of all the processes. These will be used to determine which events can be committed and removed. It also contains the receive events from messages send by the receiver of this COMM message that have been committed by the sender of the COMM message in addition to this the amount of these receive events in the message. This message will be sent by the initiator of the commit procedure only.
- **COMM4:** This message will contain the receive events from messages send by the receiver of this COMM message that have been committed by the sender of the COMM message in addition to this the amount of these receive events in the message. This message will be sent by the non-initiators of the commit procedure only.

The event log consists of an array of structs that contain the send and receive events that are stored, message size, the ID of the message that send the message, the time vector at the moment of sending, the time vector at the moment of receiving and the failure vector at the moment of sending. it also contains the index of the next checkpoint and a boolean indicating if the messages is a send or a receive event.

The checkpointing file in stable storage is organized as follows: at the start of the file an integers will be stored with the amount of total checkpoints in the file. Afterwards the checkpoints will

be stored. the header of the checkpoint will consist of two integers and the time vector of the checkpoint. The initial integers will consist of the messages stored in total at the time the checkpoint will be made and the messages stored in total at the checkpoint before this one. The actual checkpoint data will be the serialized `msg_log_t`.

The committal file will be populated with the serialized `msg_log_t` of the committed messages.

4.3 Experimental Setup

The experimental setup that was used to measure the overhead cost of the implementation operates as follows:

A parent process creates pipes for its child processes to communicate and then forks them. The child processes will then randomly send each other messages using the State class. The parent process is then able to interrupt one of its children by sending them a SIGINT interrupt. This interrupt is then caught by the child process which will then complete its current select loop and will then exit. After the child has exited the parent will restart the process by forking it again. A process will transmit its new file descriptors over UNIX domain sockets. The child processes will use a select loop to deal with the incoming messages of the other processes.

The experiments will be run on a dell XPS 13 using a intel i7-8565U CPU running at 1.80GHz, the operating system used was Arch-Linux 2024.05.01 release with the included kernel version 6.8.8. Both the implementation of the protocol and the experimental framework is written in C++ version 20.

The measurements are taken using the `chrono::high_resolution_clock` from `time.h`. The measurements of the overhead cost for sending and receiving and checkpointing are taken by only measuring how long it took to complete the `send_msg`, `recv_msg` and `checkpoint` methods. For the restart time cost measurements are taken by measuring the time it took for the class to be initialized and transmit all control messages and thus resume normal execution using increasing number of messages stored in stable storage. For the rollback time cost measurements is the time it takes from the moment the process receives the control message until it has handled the rollback procedure and resumed normal operation, each subsequent time using increasing number of messages stored in its event log to determine what effect this has on the rollback pro-

cedure. The communication overhead and stable storage memory overhead is calculated using the serialized size of the message and event log. The RAM memory overhead is calculated by the used memory space by the event log struct and the size of the accompanying vectors. The measurements are saved in a CSV file to be able to analyses the results more easily.

The implementation of the experimental setup can be found at GitHub [9].

In addition to the performance test a correctness test will be performed, various scenarios will be checked to indicate the correctness of the implementation.

5 Results and Analysis

5.0.1 Memory overhead

The implementation of the improved Peterson-Kearns protocol both for storing messages in stable storage and in communication utilizes serialized data. The overhead for storing messages in stable storage:

$$3 * sizeof(int) + 3 * sizeof(int) * n \text{ bytes}$$

the communications overhead:

$$3 * sizeof(int) + 2 * sizeof(int) * n \text{ bytes}$$

and the overhead for storing messages in random access memory:

$$96 + 3 * sizeof(int) * n \text{ bytes}$$

where

$$n = \text{number of processes}$$

This shows that when this implementation is used by a large amount of processes that the memory overhead will increase substantially. This overhead could be diminished if an `int16_t` would be used to store the time and failure vectors, however the downside to this would be that the implementation would be able to handle less messages.

5.1 Experimental Findings

5.1.1 Sending performance

The two figures 3 , 4 show the 200 moving average of time cost of sending a message. The sending cost is tested using different amounts of processes communicating. In the first figure the processes were committing and removing events

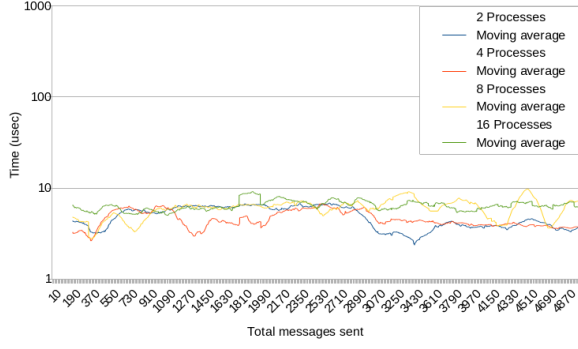


Figure 3: Sending cost for one message : commit procedure (200 MA)



Figure 4: Sending cost for one message (200 MA)

as well as checkpointing the events, in the second figure the processes were only checkpointing the events. This shows that the sending performance is not affected by the increase in processes communicating. It also shows that there is not a big difference in sending cost between whether the messages will be committed or not, however this is to be expected since the operations used for sending a message do not change between the two versions of the protocol.

5.1.2 Receiving performance




Figure 5: Receiving cost for one message : commit procedure

The two figures 5, 6 above show the receiving cost of receiving a message. The receiving cost is tested using different amounts of processes com-




Figure 6: Receiving cost for one message (200MA)

municating. The first figure shows the performance while the processes are committing and removing events. for clarity only the plot for 2 communicating processes is shown. This shows that there are numerous spikes in the time cost if we look at the second figure we can see why this occurs. In the second figure a 200 moving average is used. Here we can see that we do not have such spikes in time cost and can thus conclude that the committing and removing events is a costly procedure. we can also see that the receiving performance is not affected by the increase in processes communicating.

5.1.3 Restart performance




Figure 7: Average restart cost

The figure 7 above shows the time cost of the restart procedure. The blue line shows the time it took from the moment the process called the constructor of the State class until normal operation was resumed, the red line the amount of messages in stable memory that were stored by the process performing the restart. This shows that the time cost of the restart procedure is proportionate to the amount of messages that are stored in stable storage, furthermore it shows that restarting is relatively a cheap procedure.

5.1.4 Rollback performance

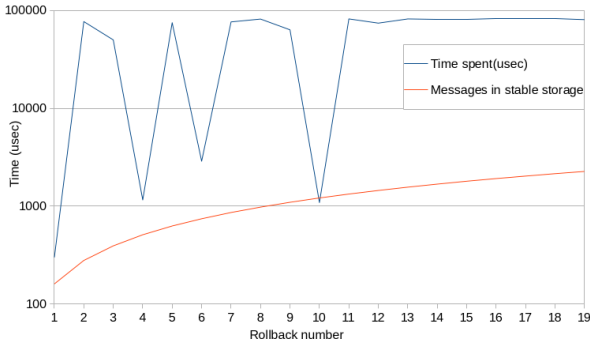


Figure 8: Rollback cost




Figure 9: Average rollback cost

The figures 8, 9 above shows the time cost of the rollback procedure. The line in blue shows the time it took from the moment the process received the control message until normal operation was resumed, the red line shows the amount of messages in volatile memory that were stored by the process performing the rollback. This shows that this is a costly procedure since it takes almost 0,1 second most of the time. However at the 1st, 4th, 6th and 10th time the rollback was performed shows a discrepancy, it takes considerably less time to complete the rollback. This is due to that there has not been a checkpoint in between the other process crashing and the process performing the rollback to receive the control message and thus it will not have to remove any checkpoints. The only thing the rollback then does is to replay, resend and remove messages from volatile memory. This indicates that the removal of checkpoints is a costly operation. In figure 10 the rollback cost when dealing with concurrent failures, all odd number rollbacks are the from the initial failure the even rollback numbers indicate the concurrent failure after the initial one.

This shows that the concurrent failure will cause the rollback to be performed much faster

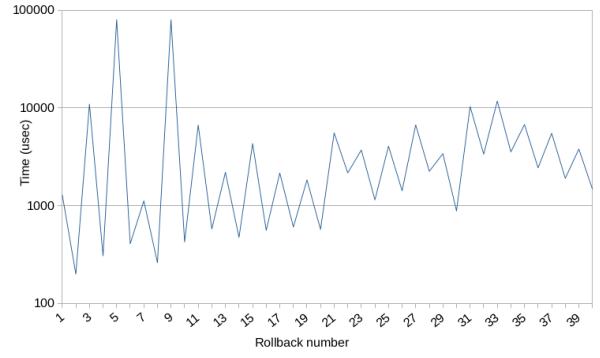


Figure 10: Concurrent rollback cost

than the rollback performed after the initial failure. This is logical since the rollback procedure has already been performed, only message that needed to be resent will again need to be resent, the other parts of the rollback have already been performed at the rollback caused by the initial failure.

5.2 Correctness indication

To provide a Correctness indication a couple of scenarios will be tested:

- *Duplicate*: Duplicate messages will be handled by checking if a set contains the time and failure vectors, afterwards the vectors will be inserted into the set. These will be removed when checkpoints get removed and no more duplicate messages can be sent and is thus not needed any more. The implementation for this is provided in appendix A:93 and A:1372
- *Duplicate-control*: Duplicate control messages will be handled by checking if a set contains its id, failure_number and restart_time, afterwards these values will be added to the set as an entry. The implementation for this is provided in appendix A:106.
- *Orphan-arrived*: Whenever a rollback procedure is called orphaned events will be removed from the message log. An event e_i is orphaned if $e_i.T^i > restart_time$. if the event e_i is a receive event then $e_i.T_{sender}^i > restart_time \wedge e_i.F_{sender}^i < failure_number$. Restart_time and failure_number are the values contained in the control message. The implementation for this is provided in appendix A:880.
- *Orphan-in-transit*: Whenever a message is received the check_orphan method is called.

This method will check if there exists a entry(j, n, T^j) in the failure log that has $msg.F^j < n \wedge msg.T^j > T^j$, if this exists the message is a orphan-in-transit message and will be ignored. the method will be provided in appendix A:129.

- *Childless-lost*: To resolve childless-lost event all send events that are not lost or orphaned but have their receive event lost in the failure will be re-transmitted. The implementation for this is provided in appendix A:966
- *Childless-Orphaned*: To resolve childless-orphan events all send events that are not lost or orphaned and have their corresponding receive event is orphaned in the rollback procedure will be replayed. The implementation for this is provided in appendix A:931.

This provides a rudimentary correctness indication for the implementation.

5.3 Future Work

For future implementation work the removal of checkpoints could be improved since it now takes a considerable performance penalty to evoke the method. Another point that could be improved is the return value when a process receives a control message and performs the rollback, right now it is just indicated that a rollback has been performed and not what messages have been removed during the rollback. Another improvement that could be made is a more specified error handling with setting the errno correctly. Some changes could also be made to add more constructors to the `msg_log_t` like the copy and assignment constructors. The serialization of the data to be send and stored could also be handled by some library like Protobuf [10] to facilitate streamlined compatibility with other programming languages. The way committed messages are stored could be altered to a more structured state to provide processes using the committed data an easy understanding of the way the data is stored. Another change that would be welcome is the addition of shadow paging this would increase the fault tolerance of the implementation to ensure data is still valid if a crash occurred in the middle of a write operation.

6 Conclusion

With this we can conclude that the memory overhead does increase proportionate to the number of processes communicating, however we can also conclude that the sending and receiving performance is not greatly affected by the number of processes communicating. The rollback performance as well as the commit receive performance shows that the removing of checkpoints a costly operation. In general is rolling back a costly operation while restarting is a relatively cheap operation.

A pet-kea.cpp

```
1  #include "pet-kea.hpp"
2
3  using namespace std;
4  typedef unordered_set<vector<int>, vector_hash> uset;
5
6  void get_fail_filename(int process_nr, char fail_filename[32])
7  {
8      sprintf(fail_filename, "fail_v_process_%d.dat", process_nr);
9  }
10
11 void get_msg_filename(int process_nr, char msg_filename[32])
12 {
13     sprintf(msg_filename, "msg_process_%d.dat", process_nr);
14 }
15
16 void get_comm_filename(int process_nr, char commit_filename[32])
17 {
18     sprintf(commit_filename, "commit_process%d.dat", process_nr);
19 }
20
21 void get_state_filename(int process_nr, char state_filename[32])
22 {
23     sprintf(state_filename, "state_process%d.dat", process_nr);
24 }
25
26 void Pet_kea::print_msg(struct msg_t *msg)
27 {
28     cout << "MSG_time_vector(";
29     for (int i = 0; i < (int)msg->time_v.size(); i++)
30     {
31         cout << msg->time_v[i];
32         if (i < (int)msg->time_v.size())
33         {
34             cout << ", ";
35         }
36     }
37     cout << ")_fail_vector(";
38     for (int i = 0; i < (int)msg->fail_v.size(); i++)
39     {
40         cout << msg->fail_v[i];
41         if (i < (int)msg->fail_v.size())
42         {
43             cout << ", ";
44         }
45     }
46     cout << ")" << endl;
47 }
48
49 void Pet_kea::print_ctrl_msg(struct ctrl_msg_t *msg)
50 {
51     cout << "CTRL_with_log(" << msg->log_entry.id << ", " << msg->log_entry.fail_nr << ", ";
52     cout << msg->log_entry.res_time << ")_messages_recieved:" << msg->recieved_cnt << endl;
53     for (set<pair<int, vector<int>>>::iterator ptr = msg->recieved_msgs.begin();
54          ptr != msg->recieved_msgs.end(); ptr++)
55     {
56         cout << "Tj:" << ptr->first << "_fail_v:";
57         for (int j = 0; j < (int)ptr->second.size(); j++)
58         {
59             cout << ptr->second[j] << ";";
60         }
61         cout << "_length:" << ptr->second.size() << endl;
62     }
63 }
64
65 int *Pet_kea::State::next_checkpoint(int *ptr)
66 {
67     int to_skip = *ptr - *(ptr + 1);
68     ptr += 2 + time_v.size();
69
70     for (int i = 0; i < to_skip; i++)
71     {
72         ptr += ((SER_LOG_SIZE + *ptr) / sizeof(int));
73     }
74     return ptr;
75 }
76
77 char *Pet_kea::State::get_msg(int i)
78 {
79     return msg_log[i].msg_buf;
80 }
81
82 char **Pet_kea::State::get_msg_log()
83 {
84     char **output_log = (char **)malloc(msg_cnt * sizeof(char *));
85     for (int i = 0; i < msg_cnt; i++)
86     {
87         output_log[i] = msg_log[i].msg_buf;
88     }
89     return output_log;
90 }
91
92 bool Pet_kea::State::check_duplicate(struct msg_t *msg)
93 {
94     vector<int> merged_time_fail_v;
95     merged_time_fail_v = msg->time_v;
96     merged_time_fail_v.insert(merged_time_fail_v.end(), msg->fail_v.begin(), msg->fail_v.end());
97     if (arrived_msgs.contains(merged_time_fail_v))
98         return true;
99 }
100
```

```

101     arrived_msgs.insert(merged_time_fail_v);
102     return false;
103 }
104
105 bool Pet_kea::State::check_duplicate_ctrl(struct fail_log_t log)
106 {
107     vector<int> fail_log_vector{log.id, log.fail_nr, log.res_time};
108     if (arrived_ctrl.contains(fail_log_vector))
109         return true;
110
111     arrived_ctrl.insert(fail_log_vector);
112     return false;
113 }
114
115 bool Pet_kea::State::check_duplicate_commit(struct msg_log_t *l_msg)
116 {
117     vector<int> merged_time_fail_v;
118     merged_time_fail_v = l_msg->time_v_sender;
119     merged_time_fail_v.insert(merged_time_fail_v.end(),
120                             l_msg->fail_v_sender.begin(), l_msg->fail_v_sender.end());
121     if (committed_msg_set.contains(merged_time_fail_v))
122         return true;
123
124     committed_msg_set.insert(merged_time_fail_v);
125     return false;
126 }
127
128 bool Pet_kea::State::check_orphaned(struct msg_t *msg)
129 {
130     for (int i = 0; i < (int)fail_log.size(); i++)
131     {
132         if (msg->fail_v[fail_log[i].id] < fail_log[i].fail_nr &&
133             msg->time_v[fail_log[i].id] > fail_log[i].res_time)
134             return true;
135         else
136             continue;
137     }
138     return false;
139 }
140
141 void Pet_kea::State::rem_checkpoints(vector<int> to_remove)
142 {
143     // read file and reconstruct it
144
145     vector<int> reverse_to_remove = to_remove;
146     reverse(reverse_to_remove.begin(), reverse_to_remove.end());
147     char filename[32];
148     get_msg_filename(id, filename);
149     msg_out.close();
150     ifstream msg_in(filename, ifstream::in | ifstream::binary);
151     msg_in.seekg(0, msg_in.end);
152     size_t file_size = msg_in.tellg();
153     msg_in.seekg(0, ifstream::beg);
154     char msg_file[file_size];
155     msg_in.read(msg_file, file_size);
156     msg_in.close();
157
158     char new_msg_file[file_size];
159
160     int *old_ptr = (int *)msg_file;
161     int *new_ptr = (int *)new_msg_file;
162
163     old_ptr++;
164     new_ptr++;
165
166     int checkpoint_msg_cnt = 0, checkpoint_last_ckpnt = 0;
167     int checkpoint_cnt = (int)checkpoints.size();
168     vector<int> new_checkpoints;
169     vector<std::vector<int>> new_ck_time_v;
170     new_checkpoints.push_back(0);
171     new_ck_time_v.push_back(ck_time_v[0]);
172     int *temp_ptr;
173
174     for (int i = 1; i < checkpoint_cnt; i++)
175     {
176         if (!reverse_to_remove.empty() && reverse_to_remove.back() == i)
177         {
178             reverse_to_remove.pop_back();
179             old_ptr = next_checkpoint(old_ptr);
180             continue;
181         }
182
183         temp_ptr = old_ptr;
184
185         checkpoint_last_ckpnt = checkpoint_msg_cnt;
186         checkpoint_msg_cnt += *old_ptr - *(old_ptr + 1);
187         *new_ptr = checkpoint_msg_cnt;
188         new_ptr++;
189         old_ptr++;
190         *new_ptr = checkpoint_last_ckpnt;
191         new_ptr++;
192         old_ptr++;
193
194         temp_ptr = next_checkpoint(temp_ptr);
195
196         memcpy(new_ptr, old_ptr, (temp_ptr - old_ptr) * sizeof(int));
197
198         new_checkpoints.push_back(checkpoint_msg_cnt);
199         new_ck_time_v.push_back(ck_time_v[i]);
200
201         new_ptr += (temp_ptr - old_ptr);
202
203

```

```

204     old_ptr = temp_ptr;
205 }
206
207 checkpoints.swap(new_checkpoints);
208 ck_time_v.swap(new_ck_time_v);
209
210 size_t test = (size_t)new_ptr;
211 size_t test2 = (size_t)(int *)&new_msg_file;
212
213 file_size = new_ptr - (int *)&new_msg_file;
214 file_size = test - test2;
215
216 new_ptr = (int *)&new_msg_file;
217
218 *new_ptr = checkpoints.size() - 1;
219
220 msg_out.open(filename, ofstream::out | ofstream::binary | ofstream::trunc);
221 msg_out.write(new_msg_file, file_size);
222 return;
223 }
224
225 int Pet_kea::State::next_checkpoint_after_rem(vector<int> removed_checkpoints, int curr_next_checkpoint)
226 {
227     int result = curr_next_checkpoint;
228     for (const int &i : removed_checkpoints)
229     {
230         if (i < curr_next_checkpoint)
231             result--;
232         else
233             break;
234     }
235     return result;
236 }
237
238 int Pet_kea::State::rem_log_entries(vector<int> to_remove, int final_index)
239 {
240     cout << id << "old_msg_cnt:" << final_index << "removing:" << to_remove.size() << endl;
241     msg_log_t *new_log = (msg_log_t *)calloc(MAX_LOG, sizeof(msg_log_t));
242
243     vector<int>::iterator curr = to_remove.begin();
244     int new_final_index = 0;
245     for (int i = 0; i < msg_cnt; i++)
246     {
247         if (curr != to_remove.end() && i == *curr)
248         {
249             curr++;
250             free(msg_log[i].msg_buf);
251             vector<int>().swap(msg_log[i].time_v_reciever);
252             vector<int>().swap(msg_log[i].time_v_sender);
253             vector<int>().swap(msg_log[i].fail_v_sender);
254             continue;
255         }
256         new_log[new_final_index] = msg_log[i];
257         new_final_index++;
258         free(msg_log[i].msg_buf);
259         vector<int>().swap(msg_log[i].time_v_reciever);
260         vector<int>().swap(msg_log[i].time_v_sender);
261         vector<int>().swap(msg_log[i].fail_v_sender);
262     }
263     free(msg_log);
264     msg_log = new_log;
265     cout << id << "new_msg_cnt:" << new_final_index << endl;
266     return new_final_index;
267 }
268
269 void Pet_kea::State::serialize_commit(struct comm_msg_t *msg, char *data)
270 {
271     int *q = (int *)data;
272     *q = (int)msg->msg_type;
273     q++;
274
275     *q = msg->sending_process_nr;
276     q++;
277
278     switch (msg->msg_type)
279     {
280     case COMM1:
281         *q = 0; // padding
282         break;
283
284     case COMM2:
285         *q = msg->time_v_j;
286         break;
287
288     case COMM3:
289         *q = msg->committed_cnt;
290         q++;
291         for (int i = 0; i < (int)time_v.size(); i++)
292         {
293             *q = msg->time_v_min[i];
294             q++;
295         }
296         for (uset::iterator ptr = msg->committed_msgs.begin();
297              ptr != msg->committed_msgs.end(); ptr++)
298         {
299             for (int i = 0; i < (int)time_v.size() * 2; i++)
300             {
301                 *q = ptr->at(i);
302                 q++;
303             }
304         }
305         break;
306     case COMM4:

```



```

307
308     *q = msg->committed_cnt;
309     q++;
310     for (uset::iterator ptr = msg->committed_msgs.begin();
311          ptr != msg->committed_msgs.end(); ptr++)
312     {
313         for (int i = 0; i < (int)time_v.size() * 2; i++)
314         {
315             *q = ptr->at(i);
316             q++;
317         }
318     }
319     break;
320
321
322     default:
323         break;
324 }
325
326
327 void Pet_kea::State::deserialize_commit(char *data, struct comm_msg_t *msg)
328 {
329     int *q = (int *)data;
330     msg->msg_type = (msg_type)*q;
331     q++;
332
333     msg->sending_process_nr = *q;
334     q++;
335     vector<int> temp_vec;
336     switch (msg->msg_type)
337     {
338     case COMM1:
339         break;
340
341     case COMM2:
342         msg->time_v_j = *q;
343         break;
344
345     case COMM3:
346
347         msg->committed_cnt = *q;
348         q++;
349         for (int i = 0; i < (int)time_v.size(); i++)
350         {
351             msg->time_v_min.push_back(*q);
352             q++;
353         }
354
355         for (int i = 0; i < msg->committed_cnt; i++)
356         {
357             for (int j = 0; j < (int)time_v.size() * 2; j++)
358             {
359                 temp_vec.push_back(*q);
360                 q++;
361             }
362
363             msg->committed_msgs.insert(temp_vec);
364             temp_vec.clear();
365         }
366         break;
367     case COMM4:
368
369         msg->committed_cnt = *q;
370         q++;
371         for (int i = 0; i < msg->committed_cnt; i++)
372         {
373             for (int j = 0; j < (int)time_v.size() * 2; j++)
374             {
375                 temp_vec.push_back(*q);
376                 q++;
377             }
378
379             msg->committed_msgs.insert(temp_vec);
380             temp_vec.clear();
381         }
382
383         break;
384
385     default:
386         break;
387 }
388
389
390 void Pet_kea::State::serialize_ctrl(struct ctrl_msg_t *msg, char *data)
391 {
392     int *q = (int *)data;
393     *q = (int)msg->msg_type;
394     q++;
395
396     *q = msg->recieved_cnt;
397     q++;
398
399     *q = msg->sending_process_nr;
400     q++;
401
402     *q = msg->log_entry.id;
403     q++;
404     *q = msg->log_entry.fail_nr;
405     q++;
406     *q = msg->log_entry.res_time;
407     q++;
408
409     for (set<pair<int, vector<int>>>::iterator ptr = msg->recieved_msgs.begin();

```

```

410         ptr != msg->recieved_msgs.end(); ptr++)
411     {
412         *q = ptr->first;
413         q++;
414         for (int j = 0; j < (int)fail_v.size(); j++)
415         {
416             *q = ptr->second[j];
417             q++;
418         }
419     }
420 }
421
422 void Pet_kea::State::deserialize_ctrl(char *data, struct ctrl_msg_t *msg)
423 {
424     int *q = (int *)data;
425     msg->msg_type = (msg_type)*q;
426     q++;
427
428     msg->recieved_cnt = *q;
429     q++;
430
431     msg->sending_process_nr = *q;
432     q++;
433
434     msg->log_entry.id = *q;
435     q++;
436     msg->log_entry.fail_nr = *q;
437     q++;
438     msg->log_entry.res_time = *q;
439     q++;
440
441     pair<int, vector<int>>> temp_pair;
442     for (int i = 0; i < msg->recieved_cnt; i++)
443     {
444
445         temp_pair.first = *q;
446         q++;
447         for (int j = 0; j < (int)fail_v.size(); j++)
448         {
449             temp_pair.second.push_back(*q);
450             q++;
451         }
452
453         msg->recieved_msgs.insert(temp_pair);
454         temp_pair.second.clear();
455     }
456 }
457
458 void Pet_kea::State::serialize(struct msg_t *msg, char *data)
459 {
460     int *q = (int *)data;
461     *q = msg->msg_type;
462     q++;
463
464     *q = msg->msg_size;
465     q++;
466
467     *q = msg->sending_process_nr;
468     q++;
469     if (msg->msg_type == MSG)
470     {
471         for (int i = 0; i < (int)time_v.size(); i++)
472         {
473             *q = msg->time_v[i];
474             q++;
475         }
476
477         for (int i = 0; i < (int)fail_v.size(); i++)
478         {
479             *q = msg->fail_v[i];
480             q++;
481         }
482     }
483
484     memcpy(q, msg->msg_buf, msg->msg_size);
485 }
486 void Pet_kea::State::deserialize(char *data, struct msg_t *msg)
487 {
488     int *q = (int *)data;
489     msg->msg_type = (msg_type)*q;
490     q++;
491
492     msg->msg_size = *q;
493     q++;
494
495     msg->sending_process_nr = *q;
496     q++;
497     if (msg->msg_type == MSG)
498     {
499         for (int i = 0; i < (int)time_v.size(); i++)
500         {
501             msg->time_v.push_back(*q);
502             q++;
503         }
504
505         for (int i = 0; i < (int)fail_v.size(); i++)
506         {
507             msg->fail_v.push_back(*q);
508             q++;
509         }
510     }
511
512     msg->msg_buf = (char *) malloc(msg->msg_size);

```

```

513     memcpy(msg->msg_buf, q, msg->msg_size);
514 }
515
516 void Pet_kea::State::serialize_log(struct msg_log_t *log, char *data)
517 {
518     int *q = (int *)data;
519     *q = log->msg_size;
520     q++;
521     *q = log->recipient;
522     q++;
523     *q = log->process_id;
524     q++;
525     *q = log->next_checkpoint;
526     q++;
527
528     for (int i = 0; i < (int)time_v.size(); i++)
529     {
530         *q = log->time_v_sender[i];
531         q++;
532     }
533
534     for (int i = 0; i < (int)time_v.size(); i++)
535     {
536         *q = log->time_v_reciever[i];
537         q++;
538     }
539
540     for (int i = 0; i < (int)fail_v.size(); i++)
541     {
542         *q = log->fail_v_sender[i];
543         q++;
544     }
545
546     memcpy(q, log->msg_buf, log->msg_size);
547 }
548 int Pet_kea::State::deserialize_log(char *data, struct msg_log_t *log)
549 {
550     int *q = (int *)data;
551     log->msg_size = *q;
552     q++;
553     log->recipient = *q;
554     q++;
555     log->process_id = *q;
556     q++;
557     log->next_checkpoint = *q;
558     q++;
559
560     for (int i = 0; i < (int)time_v.size(); i++)
561     {
562         log->time_v_sender.push_back(*q);
563         q++;
564     }
565
566     for (int i = 0; i < (int)time_v.size(); i++)
567     {
568         log->time_v_reciever.push_back(*q);
569         q++;
570     }
571
572     for (int i = 0; i < (int)fail_v.size(); i++)
573     {
574         log->fail_v_sender.push_back(*q);
575         q++;
576     }
577
578     log->msg_buf = (char *)malloc(log->msg_size);
579     memcpy(log->msg_buf, q, log->msg_size);
580
581     return (q - (int *)data) * sizeof(int) + log->msg_size;
582 }
583
584 void Pet_kea::State::serialize_state(char *data)
585 {
586     int *q = (int *)data;
587     *q = id;
588     q++;
589
590     *q = msg_cnt;
591     q++;
592
593     *q = arrived_ctrl.size();
594     q++;
595
596     *q = committed_msg_set.size();
597     q++;
598
599     *q = committed_recieve_events.size();
600     q++;
601
602     for (int i = 0; i < (int)time_v.size(); i++)
603     {
604         *q = time_v[i];
605         q++;
606     }
607
608     for (int i = 0; i < (int)time_v.size(); i++)
609     {
610         *q = time_v_min[i];
611         q++;
612     }
613
614     for (int i = 0; i < (int)time_v.size(); i++)
615     {
616         *q = commit_v[i];
617         q++;

```

```

616     }
617
618     for (int i = 0; i < (int)time_v.size(); i++)
619     {
620         *q = remove_v[i];
621         q++;
622     }
623
624     for (uset::iterator ptr = arrived_ctrl.begin(); ptr != arrived_ctrl.end(); ptr++)
625     {
626         *q = ptr->at(0);
627         q++;
628         *q = ptr->at(1);
629         q++;
630         *q = ptr->at(2);
631         q++;
632     }
633     for (uset::iterator ptr = committed_msg_set.begin(); ptr != committed_msg_set.end(); ptr++)
634     {
635         for (int i = 0; i < (int)time_v.size() * 2; i++)
636         {
637             *q = ptr->at(i);
638             q++;
639         }
640     }
641     for (uset::iterator ptr = committed_recieve_events.begin(); ptr != committed_recieve_events.end(); ptr++)
642     {
643         for (int i = 0; i < (int)time_v.size() * 2; i++)
644         {
645             *q = ptr->at(i);
646             q++;
647         }
648     }
649 }
650
651 void Pet_kea::State::deserialize_state(char *data)
652 {
653     int *q = (int *)data;
654     id = *q;
655     q++;
656
657     msg_cnt = *q;
658     q++;
659
660     int arrived_ctrl_size = *q;
661     q++;
662
663     int committed_msg_set_size = *q;
664     q++;
665
666     int committed_recieve_events_size = *q;
667     q++;
668
669     for (int i = 0; i < (int)time_v.size(); i++)
670     {
671         time_v[i] = *q;
672         q++;
673     }
674
675     for (int i = 0; i < (int)time_v.size(); i++)
676     {
677         time_v_min[i] = *q;
678         q++;
679     }
680
681     for (int i = 0; i < (int)time_v.size(); i++)
682     {
683         commit_v[i] = *q;
684         q++;
685     }
686
687     for (int i = 0; i < (int)time_v.size(); i++)
688     {
689         remove_v[i] = *q;
690         q++;
691     }
692
693     vector<int> temp_vec;
694     for (int i = 0; i < arrived_ctrl_size; i++)
695     {
696         temp_vec.push_back(*q);
697         q++;
698         temp_vec.push_back(*q);
699         q++;
700         temp_vec.push_back(*q);
701         q++;
702         arrived_ctrl.insert(temp_vec);
703         temp_vec.clear();
704     }
705
706     for (int i = 0; i < committed_msg_set_size; i++)
707     {
708         for (int j = 0; j < (int)time_v.size() * 2; j++)
709         {
710             temp_vec.push_back(*q);
711             q++;
712         }
713         committed_msg_set.insert(temp_vec);
714         temp_vec.clear();
715     }
716
717     for (int i = 0; i < committed_recieve_events_size; i++)
718     {

```

```

719     for (int j = 0; j < (int)time_v.size() * 2; j++)
720     {
721         temp_vec.push_back(*q);
722         q++;
723     }
724     committed_recieve_events.insert(temp_vec);
725     temp_vec.clear();
726 }
727 }
728
729 int Pet_kea::State::send_ctrl()
730 {
731     set<pair<int, std::vector<int>>> recvd_msgs[time_v.size()];
732
733     vector<int> cnt(time_v.size(), 0);
734     struct fail_log_t fail_log = {id, fail_v[id], time_v[id]};
735
736     pair<int, vector<int>> temp_pair;
737     for (int i = 0; i < msg_cnt; i++)
738     {
739         if (!msg_log[i].recipient)
740         {
741             continue;
742         }
743
744         temp_pair.first = msg_log[i].time_v_sender[msg_log[i].process_id];
745         temp_pair.second = msg_log[i].fail_v_sender;
746         recvd_msgs[msg_log[i].process_id].insert(temp_pair);
747         cnt[msg_log[i].process_id]++;
748     }
749
750     for (int i = 0; i < (int)time_v.size(); i++)
751     {
752         // prepare and send the ctrl messages
753         if (i == id)
754             continue;
755         struct ctrl_msg_t msg;
756         msg.msg_type = CTRL;
757         msg.sending_process_nr = id;
758         msg.log_entry = fail_log;
759         msg.recieved_cnt = cnt[i];
760         msg.recieved_msgs = recvd_msgs[i];
761
762         print_ctrl_msg(&msg);
763
764         // send the control message (serialize)
765         size_t size = SER_SIZE_CTRL_MSG_T(msg.recieved_cnt, fail_v.size());
766         char *data = (char *)malloc(size);
767         serialize_ctrl(&msg, data);
768         try
769         {
770             int ret = write(fildes[i][1], data, size);
771             if (ret < 0)
772             {
773                 throw runtime_error("failed_to_write");
774             }
775         }
776         catch (const std::exception &e)
777         {
778             std::cerr << e.what() << endl;
779             perror("write_failed");
780             return -1;
781         }
782
783         free(data);
784     }
785     return 0;
786 }
787
788 int Pet_kea::State::store_msg(struct msg_t *msg, int recipient)
789 {
790     // store the message
791     if (msg_cnt >= MAX_LOG)
792     {
793         cout << "max_log_size_reached_of_process_" << id << endl;
794         return -1;
795     }
796
797     msg_log[msg_cnt].msg_size = msg->msg_size;
798     msg_log[msg_cnt].next_checkpoint = checkpoints.size();
799     msg_log[msg_cnt].msg_buf = (char *)malloc(msg->msg_size);
800     memcpy(msg_log[msg_cnt].msg_buf, msg->msg_buf, msg->msg_size);
801     msg_log[msg_cnt].time_v_reciever = time_v;
802     if (recipient == -1)
803     {
804         msg_log[msg_cnt].time_v_sender = msg->time_v;
805         msg_log[msg_cnt].fail_v_sender = msg->fail_v;
806         msg_log[msg_cnt].process_id = msg->sending_process_nr;
807         msg_log[msg_cnt].recipient = true;
808     }
809     else
810     {
811         msg_log[msg_cnt].time_v_sender = time_v;
812         msg_log[msg_cnt].fail_v_sender = fail_v;
813         msg_log[msg_cnt].process_id = recipient;
814         msg_log[msg_cnt].recipient = false;
815     }
816
817     msg_cnt++;
818     if (SAVE_CNT != 0 && msg_cnt % SAVE_CNT == 0)
819         checkpoint();
820
821     return 0;

```

```

822 }
823
824 int Pet_kea::State::commit_msgs(vector<int> msgs)
825 {
826     char filename[32];
827     get_comm_filename(id, filename);
828
829     for (int iterator : msgs)
830     {
831         char data[msg_log[iterator].msg_size + SER_LOG_SIZE];
832         serialize_log(&msg_log[iterator], data);
833         commit_out.write(data, msg_log[iterator].msg_size + SER_LOG_SIZE);
834     }
835     commit_out.flush();
836
837     return 0;
838 }
839
840 int Pet_kea::State::rollback(struct ctrl_msg_t *msg)
841 {
842     cout << "entered_the_rollback_section" << endl;
843     print_ctrl_msg(msg);
844
845     // RB.2
846     char filename[32];
847     get_fail_filename(id, filename);
848     ofstream fail_out(filename, ofstream::out | ofstream::binary | ofstream::app);
849
850     struct fail_log_t entry = {msg->sending_process_nr, msg->log_entry.fail_nr, msg->log_entry.res_time};
851     fail_out.write((char *)&entry, sizeof(fail_log_t));
852     fail_out.close();
853
854     fail_log.push_back(entry);
855
856     // RB.2.3
857     fail_v[msg->sending_process_nr] = msg->log_entry.fail_nr;
858     if (time_v[msg->sending_process_nr] > msg->log_entry.res_time)
859     {
860         // RB.2.1 remove checkpoints T^i > crT^i, set state and T to latest checkpoint, replays
861         vector<int> checkpoints_to_remove;
862
863         for (int i = 0; i < (int)checkpoints.size(); i++)
864         {
865             if (ck_time_v[i].at(msg->sending_process_nr) > msg->log_entry.res_time)
866                 checkpoints_to_remove.push_back(i);
867         }
868
869         rem_checkpoints(checkpoints_to_remove);
870
871         // set state and time_v to latestck
872         int prev_cnt = msg_cnt;
873         msg_cnt = checkpoints.back();
874         int temp_msg_cnt = msg_cnt;
875         time_v = ck_time_v.back();
876
877         // replay messages
878         std::vector<int> indices_to_rem;
879
880         for (int i = msg_cnt; i < prev_cnt; i++)
881         {
882             if (msg_log[i].time_v_sender[msg->sending_process_nr] <= msg->log_entry.res_time &&
883                 msg_log[i].time_v_sender[id] >= ck_time_v.back()[id])
884             {
885                 // replay messages only inc time_v and msg_cnt
886                 cout << "replayed_msg" << endl;
887                 msg_cnt++;
888                 if (msg_log[i].recipient)
889                 {
890                     time_v[id]++;
891                     for (int j = 0; j < (int)time_v.size(); j++)
892                     {
893                         if (j == id)
894                             continue;
895                         time_v.at(j) = max(msg_log[i].time_v_sender[j], time_v[j]);
896                     }
897                 }
898                 else
899                 {
900                     time_v[id]++;
901                 }
902             }
903             else if (!msg_log[i].recipient &&
904                 msg_log[i].time_v_sender[msg->sending_process_nr] > msg->log_entry.res_time)
905             {
906                 // add to remove vector
907                 indices_to_rem.push_back(i);
908             }
909             else if (msg_log[i].recipient &&
910                 msg_log[i].time_v_reciever[msg->sending_process_nr] > msg->log_entry.res_time &&
911                 msg_log[i].time_v_sender[msg->sending_process_nr] > msg->log_entry.res_time &&
912                 msg_log[i].fail_v_sender[msg->sending_process_nr] < msg->log_entry.fail_nr)
913             {
914                 // add to remove vector
915                 indices_to_rem.push_back(i);
916             }
917         }
918         if (!indices_to_rem.empty())
919         {
920             msg_cnt = rem_log_entries(indices_to_rem, prev_cnt);
921             indices_to_rem.clear();
922         }
923
924         // RB.2.2

```

```

925
926 // RB.3
927 prev_cnt = msg_cnt;
928 for (int i = temp_msg_cnt; i < prev_cnt; i++)
929 {
930     // move recv event to the back
931     if (msg_log[i].recipient &&
932         msg_log[i].time_v_reciever[msg->sending_process_nr] > msg->log_entry.res_time)
933     {
934         cout << id << "moved_RECV_event_to_the_back" << endl;
935         if (msg_cnt >= MAX_LOG)
936         {
937             cout << "max_log_size_reached_of_process_" << id << endl;
938             return -1;
939         }
940
941         msg_log[msg_cnt].msg_size = msg_log[i].msg_size;
942         msg_log[msg_cnt].recipient = true;
943         msg_log[msg_cnt].process_id = msg_log[i].process_id;
944         msg_log[msg_cnt].msg_buf = (char *)malloc(msg_log[i].msg_size);
945         memcpy(msg_log[msg_cnt].msg_buf, msg_log[i].msg_buf, msg_log[i].msg_size);
946
947         msg_log[msg_cnt].time_v_sender = msg_log[i].time_v_sender;
948         msg_log[msg_cnt].fail_v_sender = msg_log[i].fail_v_sender;
949         msg_log[msg_cnt].time_v_reciever = msg_log[i].time_v_reciever;
950
951         msg_cnt++;
952
953         time_v[id]++;
954         for (int j = 0; j < (int)time_v.size(); j++)
955         {
956             if (j == id)
957                 continue;
958             time_v.at(j) = max(msg_log[i].time_v_sender[i], time_v[i]);
959         }
960
961         // remove the duplicate msg from the log
962         indices_to_rem.push_back(i);
963     }
964
965     // retransmit send events that have not arrived RB.3.3
966     if (!msg_log[i].recipient && msg_log[i].process_id == msg->sending_process_nr &&
967         !(msg->recieved_msgs.contains(pair<int, vector<int>>(msg_log[i].time_v_sender[id],
968             msg_log[i].fail_v_sender))))
969     {
970         cout << id << "retransmitted_msg_Tj:" << msg_log[i].time_v_sender[id] << "fail_v:";
971         for (int j = 0; j < (int)fail_v.size(); j++)
972         {
973             cout << msg_log[i].fail_v_sender[j] << ";";
974         }
975
976         cout << "res_time:" << msg_log[i].time_v_sender[msg->sending_process_nr] << endl;
977         struct msg_t retransmit_msg;
978         retransmit_msg.msg_type = MSG;
979         retransmit_msg.sending_process_nr = id;
980         retransmit_msg.time_v = msg_log[i].time_v_sender;
981         retransmit_msg.fail_v = msg_log[i].fail_v_sender;
982         retransmit_msg.msg_size = msg_log[i].msg_size;
983         retransmit_msg.msg_buf = (char *)malloc(retransmit_msg.msg_size * sizeof(char));
984         memcpy(retransmit_msg.msg_buf, msg_log[i].msg_buf, retransmit_msg.msg_size);
985
986         char data[SER_MSG_SIZE + msg_log[i].msg_size];
987         serialize(&retransmit_msg, data);
988
989         // send the message
990         try
991         {
992             if (write(fildes[msg_log[i].process_id][1], data, SER_MSG_SIZE + msg_log[i].msg_size) <
993                 0)
994                 throw runtime_error("failed_to_write");
995         }
996         catch (const std::exception &e)
997         {
998             std::cerr << e.what() << endl;
999             perror("write_failed");
1000             return -1;
1001         }
1002     }
1003     if (!indices_to_rem.empty())
1004     {
1005
1006         msg_cnt = rem_log_entries(indices_to_rem, msg_cnt);
1007     }
1008
1009     checkpoint();
1010 }
1011 return 0;
1012 }
1013
1014 Pet_kea::State::State(int process_nr, int process_cnt, int (*fd)[2], bool restart)
1015 : id(process_nr),
1016   time_v(process_cnt, 0),
1017   fail_v(process_cnt, 0),
1018   msg_cnt(0),
1019   commit_v(process_cnt, false),
1020   remove_v(process_cnt, false),
1021   arrived_msgs()
1022 {
1023
1024     time_v_min = vector(process_cnt, 0);
1025     fildes = (int **)malloc(process_cnt * sizeof(int *));
1026     for (int i = 0; i < process_cnt; i++)

```

```

1027 {
1028     fildes[i] = (int *) malloc(2 * sizeof(int));
1029     fildes[i][0] = fd[i][0];
1030     fildes[i][1] = fd[i][1];
1031 }
1032 if (restart)
1033 {
1034     char filename[32];
1035     get_state_filename(id, filename);
1036     ifstream state_in(filename, ifstream::in | ifstream::binary);
1037     state_in.seekg(0, ifstream::end);
1038     size_t file_size = state_in.tellg();
1039     state_in.seekg(0, ifstream::beg);
1040     char state_file[file_size];
1041     state_in.read(state_file, file_size);
1042     state_in.close();
1043     deserialize_state(state_file);
1044
1045     get_msg_filename(id, filename);
1046     ifstream msg_in(filename, ifstream::in | ifstream::binary);
1047     msg_in.seekg(0, msg_in.end);
1048     file_size = msg_in.tellg();
1049     msg_in.seekg(0, ifstream::beg);
1050     char msg_file[file_size];
1051     msg_in.read(msg_file, file_size);
1052     msg_in.close();
1053
1054     int *curr_pos = (int *) msg_file;
1055
1056     int num_checkpoints = *curr_pos;
1057     curr_pos++;
1058
1059     msg_log = (msg_log_t *) calloc(MAX_LOG, sizeof(msg_log_t));
1060
1061     checkpoints.push_back(0);
1062     ck_time_v.push_back(vector<int>(process_cnt, 0));
1063
1064     int read_msg_cnt = 0;
1065     for (int i = 0; i < num_checkpoints; i++)
1066     {
1067         int ck_msg_cnt = *curr_pos;
1068         checkpoints.push_back(*curr_pos);
1069         curr_pos++;
1070         int to_read = ck_msg_cnt - *curr_pos;
1071         curr_pos++;
1072         std::vector<int> temp_ck_time_v;
1073
1074         for (int j = 0; j < (int)time_v.size(); j++)
1075         {
1076             temp_ck_time_v.push_back(*curr_pos);
1077             curr_pos++;
1078         }
1079         ck_time_v.push_back(temp_ck_time_v);
1080
1081         for (int k = 0, ret = 0; k < to_read; k++, read_msg_cnt++)
1082         {
1083             ret = deserialize_log((char *)curr_pos, &msg_log[read_msg_cnt]);
1084             curr_pos += ret / sizeof(int);
1085         }
1086     }
1087
1088     // detect lost messages if crash happened during checkpoint
1089     if (msg_cnt != read_msg_cnt)
1090     {
1091         msg_cnt = read_msg_cnt;
1092         if (msg_cnt == 0)
1093         {
1094             time_v = vector<int>(process_cnt, 0);
1095         }
1096         else
1097         {
1098             time_v = msg_log[msg_cnt - 1].time_v_sender;
1099         }
1100     }
1101
1102     // insert the arrived messages into arrived_msgs
1103     vector<int> merged_time_fail_v;
1104
1105     for (int i = 0; i < msg_cnt; i++)
1106     {
1107         if (msg_log->recipient)
1108         {
1109             merged_time_fail_v = msg_log[i].time_v_sender;
1110             merged_time_fail_v.insert(merged_time_fail_v.end(),
1111                                     msg_log[i].fail_v_sender.begin(), msg_log[i].fail_v_sender.end());
1112             arrived_msgs.insert(merged_time_fail_v);
1113         }
1114     }
1115
1116     msg_out.open(filename, ofstream::out | ofstream::binary | ofstream::ate | ofstream::in);
1117     get_comm_filename(id, filename);
1118     commit_out.open(filename, ofstream::out | ofstream::binary | ofstream::app);
1119     get_fail_filename(id, filename);
1120
1121     ifstream fail_in(filename, ifstream::in | ifstream::binary);
1122
1123     fail_in.seekg(0, fail_in.end);
1124     file_size = fail_in.tellg();
1125     fail_in.seekg(0, ifstream::beg);
1126     char fail_file[file_size];
1127     fail_in.read(fail_file, file_size);
1128     fail_in.close();
1129     curr_pos = (int *) fail_file;

```



```

1130
1131 fail_log_t temp_fail_log;
1132 while (curr_pos < (int *) (fail_file + file_size))
1133 {
1134     temp_fail_log.id = *curr_pos;
1135     curr_pos++;
1136     temp_fail_log.fail_nr = *curr_pos;
1137     curr_pos++;
1138     temp_fail_log.res_time = *curr_pos;
1139     curr_pos++;
1140     fail_log.push_back(temp_fail_log);
1141     fail_v[temp_fail_log.id] = max(fail_v[temp_fail_log.id], temp_fail_log.fail_nr);
1142 }
1143
1144 ofstream fail_out(filename, ofstream::out | ofstream::binary | ofstream::app);
1145 fail_out.seekp(0, ofstream::end);
1146 fail_v[id]++;
1147 struct fail_log_t entry = {id, fail_v[id], time_v[id]};
1148 fail_out.write((char *)&entry, sizeof(fail_log_t));
1149 fail_out.close();
1150
1151 fail_log.push_back(entry);
1152
1153 send_ctrl();
1154 }
1155 else
1156 {
1157     char filename[32];
1158     get_fail_filename(id, filename);
1159     ofstream fail_out(filename, ofstream::out | ofstream::binary | ofstream::trunc);
1160     get_msg_filename(id, filename);
1161     msg_out.open(filename, ofstream::out | ofstream::binary | ofstream::trunc);
1162     get_comm_filename(id, filename);
1163     commit_out.open(filename, ofstream::out | ofstream::binary | ofstream::trunc);
1164
1165     struct fail_log_t entry = {id, 0, 0};
1166     fail_out.write((char *)&entry, sizeof(fail_log_t));
1167     fail_out.close();
1168
1169     fail_log.push_back(fail_log_t(id, 0, 0));
1170
1171     msg_log = (msg_log_t *) calloc(MAX_LOG, sizeof(msg_log_t));
1172     checkpoints.push_back(0);
1173     ck_time_v.push_back(vector<int>(process_cnt, 0));
1174 }
1175 SAVE_CNT = 0;
1176 }
1177
1178 Pet_kea::State::~~State()
1179 {
1180     for (int i = msg_cnt - 1; i >= 0; i--)
1181     {
1182         std::vector<int>().swap(msg_log[i].time_v_reciever);
1183         std::vector<int>().swap(msg_log[i].time_v_sender);
1184         std::vector<int>().swap(msg_log[i].fail_v_sender);
1185
1186         free(msg_log[i].msg_buf);
1187     }
1188
1189     free(msg_log);
1190     for (int i = 0; i < (int)time_v.size(); i++)
1191     {
1192         free(fildes[i]);
1193     }
1194
1195     free(fildes);
1196     msg_out.close();
1197     commit_out.close();
1198 }
1199
1200 int Pet_kea::State::checkpoint()
1201 {
1202     // write state and time vector at the start of the file
1203
1204     int *update = (int *) malloc(sizeof(int) * 2);
1205     *update = msg_cnt;
1206     update++;
1207     *update = checkpoints.back();
1208     update--;
1209
1210     msg_out.seekp(0, ofstream::beg);
1211     int *num_checkpoints = (int *) malloc(sizeof(int));
1212     *num_checkpoints = checkpoints.size();
1213     msg_out.write((char *) num_checkpoints, sizeof(int));
1214     free(num_checkpoints);
1215
1216     int *time_v_buffer = (int *) malloc(sizeof(int) * time_v.size());
1217     for (int i = 0; i < (int)time_v.size(); i++)
1218     {
1219         time_v_buffer[i] = time_v[i];
1220     }
1221
1222     // append last messages
1223     msg_out.seekp(0, ofstream::end);
1224
1225     msg_out.write((char *) update, sizeof(int) * 2);
1226     free(update);
1227
1228     msg_out.write((char *) time_v_buffer, sizeof(int) * time_v.size());
1229
1230     free(time_v_buffer);
1231
1232     for (int i = checkpoints.back(); i < msg_cnt; i++)

```

```

1233 {
1234     char data[msg_log[i].msg_size + SER_LOG_SIZE];
1235     serialize_log(&msg_log[i], data);
1236     msg_out.write(data, msg_log[i].msg_size + SER_LOG_SIZE);
1237 }
1238 checkpoints.push_back(msg_cnt);
1239 ck_time_v.push_back(time_v);
1240 msg_out.flush();
1241
1242 char filename[32];
1243 get_state_filename(id, filename);
1244
1245 ofstream state_out(filename, ofstream::out | ofstream::binary | ofstream::trunc);
1246 int state_size = SER_STATE_SIZE(arrived_ctrl.size(),
1247                                 committed_msg_set.size(), committed_recieve_events.size());
1248
1249 char data[state_size];
1250 serialize_state(data);
1251
1252 state_out.write(data, state_size);
1253 state_out.close();
1254 return 0;
1255 }
1256
1257 int Pet_kea::send_void(char *input, int fildes[2], int size)
1258 {
1259     char data[SER_VOID_SIZE + size];
1260
1261     int *q = (int *)data;
1262     *q = VOID;
1263     q++;
1264     *q = size;
1265     q++;
1266     memcpy(q, input, size);
1267
1268     try
1269     {
1270         if (write(fildes[1], data, SER_VOID_SIZE + size) < 0)
1271             throw runtime_error("failed_to_write");
1272     }
1273     catch (const std::exception &e)
1274     {
1275         std::cerr << e.what() << endl;
1276         perror("write_failed");
1277         return -1;
1278     }
1279     return 0;
1280 }
1281
1282 int Pet_kea::State::signal_commit()
1283 {
1284     struct comm_msg_t msg;
1285     msg.msg_type = COMM1;
1286     msg.sending_process_nr = id;
1287     char data[SER_COMM1_SIZE];
1288     serialize_commit(&msg, data);
1289
1290     // write to all other processes
1291     try
1292     {
1293         for (int i = 0; i < (int)time_v.size(); i++)
1294         {
1295             if (i == id)
1296                 continue;
1297             if (write(fildes[i][1], data, SER_COMM1_SIZE) < 0)
1298                 throw runtime_error("failed_to_write");
1299         }
1300     }
1301     catch (const std::exception &e)
1302     {
1303         std::cerr << e.what() << endl;
1304         perror("write_failed");
1305         return -1;
1306     }
1307     commit_v[id] = true;
1308     return 0;
1309 }
1310
1311 int Pet_kea::State::send_commit(int target_id)
1312 {
1313     struct comm_msg_t msg;
1314     msg.msg_type = COMM2;
1315     msg.sending_process_nr = id;
1316     msg.time_v_j = ck_time_v.back().at(id);
1317
1318     char data[SER_COMM2_SIZE];
1319     serialize_commit(&msg, data);
1320     // write back to sender
1321     try
1322     {
1323         if (write(fildes[target_id][1], data, SER_COMM2_SIZE) < 0)
1324             throw runtime_error("failed_to_write");
1325     }
1326     catch (const std::exception &e)
1327     {
1328         std::cerr << e.what() << endl;
1329         perror("write_failed");
1330         return -1;
1331     }
1332     return 0;
1333 }
1334
1335 int Pet_kea::State::remove_data()

```

```

1336 {
1337     vector<int> indices_to_remove, checkpoints_to_remove;
1338     vector<int> temp_vec;
1339     unordered_set<std::vector<int>, vector_hash> next_committed_recieve_events;
1340
1341     // remove checkpoint
1342     for (int i = 1; i < (int)checkpoints.size() - 1; i++)
1343     {
1344         if (ck_time_v[i + 1] <= time_v_min)
1345         {
1346             checkpoints_to_remove.push_back(i);
1347         }
1348     }
1349
1350     // remove messages
1351     for (int i = 0; i < msg_cnt; i++)
1352     {
1353         temp_vec.clear();
1354         temp_vec = msg_log[i].time_v_sender;
1355         temp_vec.insert(temp_vec.end(), msg_log[i].fail_v_sender.begin(), msg_log[i].fail_v_sender.end());
1356         if (msg_log[i].next_checkpoint >= (int)checkpoints.size())
1357         {
1358             msg_log[i].next_checkpoint = next_checkpoint_after_rem(checkpoints_to_remove,
1359                                                                     msg_log[i].next_checkpoint);
1360
1361             if (!msg_log[i].recipient && committed_recieve_events.contains(temp_vec))
1362             {
1363                 next_committed_recieve_events.insert(temp_vec);
1364             }
1365             continue;
1366         }
1367         if (msg_log[i].recipient && ck_time_v.at(msg_log[i].next_checkpoint) <= time_v_min)
1368         {
1369             // remove from msg_log
1370             indices_to_remove.push_back(i);
1371             committed_msg_set.erase(temp_vec);
1372             if (msg_log[i].fail_v_sender[id] == fail_v[id])
1373             {
1374                 // remove the entry from the set
1375                 vector<int> merged_time_fail_v;
1376                 merged_time_fail_v = msg_log[i].time_v_sender;
1377                 merged_time_fail_v.insert(merged_time_fail_v.end(),
1378                                         msg_log[i].fail_v_sender.begin(), msg_log[i].fail_v_sender.end());
1379                 arrived_msgs.erase(merged_time_fail_v);
1380             }
1381             continue;
1382         }
1383     }
1384     if (!msg_log[i].recipient && committed_recieve_events.contains(temp_vec))
1385     {
1386         if (ck_time_v.at(msg_log[i].next_checkpoint) <= time_v_min)
1387         {
1388             // remove from msg_log
1389             indices_to_remove.push_back(i);
1390             committed_msg_set.erase(temp_vec);
1391             continue;
1392         }
1393         else
1394         {
1395             next_committed_recieve_events.insert(temp_vec);
1396         }
1397     }
1398     msg_log[i].next_checkpoint = next_checkpoint_after_rem(checkpoints_to_remove,
1399                                                         msg_log[i].next_checkpoint);
1400 }
1401 msg_cnt = rem_log_entries(indices_to_remove, msg_cnt);
1402
1403 rem_checkpoints(checkpoints_to_remove);
1404 committed_recieve_events.clear();
1405 committed_recieve_events = next_committed_recieve_events;
1406
1407 // remove fail_log not possible in asynchronos setting
1408 return 0;
1409 }
1410
1411 int Pet_kea::State::commit(bool is_instigator)
1412 {
1413     // commit the messages
1414     vector<int> committed_msgs;
1415     for (int i = 0; i < msg_cnt; i++)
1416     {
1417         if ((msg_log[i].recipient ? msg_log[i].time_v_reciever : msg_log[i].time_v_sender) <= time_v_min &&
1418             !check_duplicate_commit(&msg_log[i]))
1419         {
1420             committed_msgs.push_back(i);
1421         }
1422     }
1423     commit_msgs(vector<int>(committed_msgs));
1424
1425     // remove information
1426     unordered_set<vector<int>, vector_hash> committed_set[time_v.size()];
1427     struct comm_msg_t msg;
1428     if (is_instigator)
1429     {
1430         msg.msg_type = COMM3;
1431         msg.time_v_min = time_v_min;
1432     }
1433     else
1434     {
1435         msg.msg_type = COMM4;
1436     }
1437 }
1438

```

```

1439 msg.sending_process_nr = id;
1440 vector<int> merged_vector;
1441
1442 int it;
1443 while (!committed_msgs.empty())
1444 {
1445     it = committed_msgs.back();
1446     if (msg_log[it].recipient)
1447     {
1448         merged_vector = msg_log[it].time_v_sender;
1449         merged_vector.insert(merged_vector.end(),
1450                             msg_log[it].fail_v_sender.begin(), msg_log[it].fail_v_sender.end());
1451         committed_set[msg_log[it].process_id].insert(merged_vector);
1452         merged_vector.clear();
1453     }
1454
1455     committed_msgs.pop_back();
1456 }
1457
1458 remove_v[id] = true;
1459 int max_cnt = 0;
1460 for (int i = 0; i < (int)time_v.size(); i++)
1461 {
1462     if (max_cnt < (int)committed_set[i].size())
1463         max_cnt = committed_set[i].size();
1464 }
1465 char *data = (char *)malloc(msg.msg_type == COMM3 ? SER_COMM3_SIZE(max_cnt) : SER_COMM4_SIZE(max_cnt));
1466 // write to all other processes
1467 for (int i = 0; i < (int)time_v.size(); i++)
1468 {
1469     if (i == id)
1470         continue;
1471
1472     msg.committed_msgs = committed_set[i];
1473     msg.committed_cnt = committed_set[i].size();
1474
1475     if (is_instigator)
1476     {
1477         cout << id << "_sending_comm3_to_" << i << "commit_cnt:_" << msg.committed_cnt << "time_v_min:";
1478         for (int j = 0; j < (int)time_v.size(); j++)
1479         {
1480             cout << msg.time_v_min[j] << " ";
1481         }
1482         cout << endl;
1483     }
1484     else
1485     {
1486         cout << id << "_sending_comm4_to_" << i << "commit_cnt:_" << msg.committed_cnt << endl;
1487     }
1488
1489     serialize_commit(&msg, data);
1490
1491     if (write(fildes[i][1], data,
1492             (msg.msg_type == COMM3 ? SER_COMM3_SIZE(msg.committed_cnt)
1493              : SER_COMM4_SIZE(msg.committed_cnt))) < 0)
1494         throw runtime_error("failed_to_write");
1495
1496     msg.committed_msgs.clear();
1497 }
1498 free(data);
1499 return 0;
1500 }
1501
1502 int Pet_kea::State::send_msg(char *input, int process_id, int size)
1503 {
1504     // inc T~i
1505     time_v[id]++;
1506
1507     struct msg_t msg;
1508
1509     msg.msg_type = MSG;
1510     msg.sending_process_nr = id;
1511     msg.time_v = time_v;
1512     msg.fail_v = fail_v;
1513     msg.msg_size = size;
1514     msg.msg_buf = (char *)malloc(size * sizeof(char));
1515     memcpy(msg.msg_buf, input, size * sizeof(char));
1516
1517     char data[SER_MSG_SIZE + size];
1518     serialize(&msg, data);
1519
1520     // send the message
1521     int ret;
1522     try
1523     {
1524         {
1525             ret = write(fildes[process_id][1], data, SER_MSG_SIZE + size);
1526             if (ret < 0)
1527             {
1528                 throw runtime_error("failed_to_write");
1529             }
1530         }
1531     } catch (const std::exception &e)
1532     {
1533         std::cerr << e.what() << endl;
1534         perror("write_failed");
1535         return -1;
1536     }
1537
1538     store_msg(&msg, process_id);
1539
1540     return 0;
1541 }

```

```

1542
1543 int Pet_kea::State::recv_msg(int fildes[2], char *output, int size)
1544 {
1545     // read message
1546     int ret;
1547     size_t init_read_size = SER_COMM1_SIZE;
1548
1549     try
1550     {
1551         char *extra_data, *data = (char *)malloc(SER_MSG_SIZE + size * sizeof(char));
1552         extra_data = data;
1553         ret = read(fildes[0], data, init_read_size);
1554
1555         if (ret < 0)
1556             throw runtime_error("failed_to_read");
1557
1558         extra_data += init_read_size;
1559
1560         int *q = (int *)data;
1561
1562         if (MSG == (msg_type)*q)
1563         {
1564             ret = read(fildes[0], extra_data, SER_MSG_SIZE + size - init_read_size);
1565             if (ret < 0)
1566                 throw runtime_error("failed_to_read");
1567
1568             struct msg_t msg;
1569             deserialize(data, &msg);
1570             free(data);
1571
1572             if (check_duplicate(&msg))
1573             {
1574                 return 3;
1575             }
1576
1577             if (check_orphaned(&msg))
1578             {
1579                 return 3;
1580             }
1581
1582             time_v[id]++;
1583             for (int i = 0; i < (int)time_v.size(); i++)
1584             {
1585                 if (i == id)
1586                     continue;
1587                 time_v.at(i) = max(msg.time_v[i], time_v[i]);
1588             }
1589
1590             // inc T^i and inc T^j to max(T^j of send event, prev event T^j)
1591
1592             store_msg(&msg, -1);
1593
1594             memcpy(output, msg.msg_buf, msg.msg_size);
1595         }
1596         else if (CTRL == (msg_type)*q)
1597         {
1598
1599             q++;
1600
1601             char *c_data = (char *)malloc(SER_SIZE_CTRL_MSG_T(*q, fail_v.size()));
1602             memcpy(c_data, data, init_read_size);
1603
1604             extra_data = c_data + init_read_size;
1605
1606             ret = read(fildes[0], extra_data, SER_SIZE_CTRL_MSG_T(*q, fail_v.size()) - init_read_size);
1607             if (ret < 0)
1608                 throw runtime_error("failed_to_read");
1609
1610             struct ctrl_msg_t c_msg;
1611             deserialize_ctrl(c_data, &c_msg);
1612             free(c_data);
1613             free(data);
1614             if (check_duplicate_ctrl(c_msg.log_entry))
1615             {
1616
1617                 for (int i = 0; i < msg_cnt; i++)
1618                 {
1619                     if (!msg_log[i].recipient && msg_log[i].process_id == c_msg.sending_process_nr &&
1620                         !(c_msg.recieved_msgs.contains(pair<int, vector<int>>(msg_log[i].time_v_sender[id],
1621                             msg_log[i].fail_v_sender))))
1622                     {
1623                         struct msg_t retransmit_msg;
1624                         retransmit_msg.msg_type = MSG;
1625                         retransmit_msg.sending_process_nr = id;
1626                         retransmit_msg.time_v = msg_log[i].time_v_sender;
1627                         retransmit_msg.fail_v = msg_log[i].fail_v_sender;
1628                         retransmit_msg.msg_size = msg_log[i].msg_size;
1629                         retransmit_msg.msg_buf = (char *)malloc(retransmit_msg.msg_size * sizeof(char));
1630                         memcpy(retransmit_msg.msg_buf, msg_log[i].msg_buf, retransmit_msg.msg_size);
1631
1632                         char data[SER_MSG_SIZE + msg_log[i].msg_size];
1633                         serialize(&retransmit_msg, data);
1634
1635                         // send the message
1636                         if (write(this->fildes[msg_log[i].process_id][1], data,
1637                             SER_MSG_SIZE + msg_log[i].msg_size) < 0)
1638                             throw runtime_error("failed_to_write");
1639                     }
1640                 }
1641
1642                 return 2;
1643             }
1644

```

```

1645         rollback(&c_msg);
1646         return 2;
1647     }
1648     else if (VOID == (msg_type)*q)
1649     {
1650         q++;
1651
1652         int v_size = *q;
1653
1654         char *v_data = (char *)malloc(SER_MSG_SIZE + v_size);
1655         memcpy(v_data, data, init_read_size);
1656         free(data);
1657         extra_data = v_data + init_read_size;
1658
1659         ret = read(fildes[0], extra_data, SER_VOID_SIZE + v_size - init_read_size);
1660         if (ret < 0)
1661             throw runtime_error("failed_to_read");
1662
1663         // process the void message aka give to method caller
1664         q = (int *)v_data;
1665         q++;
1666         q++;
1667         memcpy(output, q, v_size);
1668         free(v_data);
1669         return 1;
1670     }
1671     else if (COMM1 == (msg_type)*q)
1672     {
1673         cout << id << "_entered_COMM1" << endl;
1674         q++;
1675         send_commit(*q);
1676         free(data);
1677         return 4;
1678     }
1679     else if (COMM2 == (msg_type)*q)
1680     {
1681         cout << id << "_entered_COMM2" << endl;
1682
1683         ret = read(fildes[0], extra_data, SER_COMM2_SIZE - init_read_size);
1684         if (ret < 0)
1685             throw runtime_error("failed_to_read");
1686
1687         struct comm_msg_t comm2_msg;
1688         deserialize_commit(data, &comm2_msg);
1689         free(data);
1690         time_v_min[comm2_msg.sending_process_nr] = comm2_msg.time_v_j;
1691         commit_v[comm2_msg.sending_process_nr] = true;
1692         if (commit_v == vector<bool>(time_v.size(), true))
1693         {
1694             commit_v.flip();
1695             time_v_min[id] = ck_time_v.back().at(id);
1696             commit(true);
1697         }
1698         return 4;
1699     }
1700     else if (COMM3 == (msg_type)*q)
1701     {
1702         cout << id << "_entered_COMM3" << endl;
1703         q++;
1704         q++;
1705         char *comm3_data = (char *)malloc(SER_COMM3_SIZE(*q));
1706         memcpy(comm3_data, data, init_read_size);
1707
1708         extra_data = comm3_data + init_read_size;
1709
1710         ret = read(fildes[0], extra_data, SER_COMM3_SIZE(*q) - init_read_size);
1711         if (ret < 0)
1712             throw runtime_error("failed_to_read");
1713
1714         struct comm_msg_t comm3_msg;
1715         deserialize_commit(comm3_data, &comm3_msg);
1716         free(comm3_data);
1717         free(data);
1718         time_v_min = comm3_msg.time_v_min;
1719
1720         commit(false);
1721         committed_recieve_events.insert(comm3_msg.committed_msgs.begin(), comm3_msg.committed_msgs.end());
1722         ;
1723         remove_v[comm3_msg.sending_process_nr] = true;
1724         return 4;
1725     }
1726     else if (COMM4 == (msg_type)*q)
1727     {
1728         cout << id << "_entered_COMM4" << endl;
1729         q++;
1730         q++;
1731         char *comm4_data = (char *)malloc(SER_COMM4_SIZE(*q));
1732         memcpy(comm4_data, data, init_read_size);
1733
1734         extra_data = comm4_data + init_read_size;
1735
1736         ret = read(fildes[0], extra_data, SER_COMM4_SIZE(*q) - init_read_size);
1737         if (ret < 0)
1738             throw runtime_error("failed_to_read");
1739
1740         struct comm_msg_t comm4_msg;
1741         deserialize_commit(comm4_data, &comm4_msg);
1742         free(comm4_data);
1743         free(data);
1744         committed_recieve_events.insert(comm4_msg.committed_msgs.begin(), comm4_msg.committed_msgs.end());
1745         ;
1746         remove_v[comm4_msg.sending_process_nr] = true;
1747         if (remove_v == vector<bool>(time_v.size(), true))

```

```

1746         {
1747             remove_v.flip();
1748             remove_data();
1749         }
1750         return 4;
1751     }
1752 }
1753 catch (const std::exception &e)
1754 {
1755     std::cerr << e.what() << endl;
1756     perror("failed_in_rcv_msg");
1757 }
1758
1759 return 0;
1760 }
1761
1762 int Pet_kea::State::update_fd(int process_id, int fd[2])
1763 {
1764     if (process_id < 0 || process_id >= (int)time_v.size())
1765         return -1;
1766
1767     fildes[process_id][0] = fd[0];
1768     fildes[process_id][1] = fd[1];
1769     return 0;
1770 }

```

B pet-kea.hpp

```

1  /**
2   * @file pet-kea.hpp
3   * @brief Header file containing the state class and its member functions
4   */
5
6  #ifndef PETKEA_HPP
7  #define PETKEA_HPP
8
9  #include <unistd.h>
10 #include <vector>
11 #include <set>
12 #include <unistd.h>
13 #include <stdlib.h>
14 #include <stdio.h>
15 #include <iostream>
16 #include <fstream>
17 #include <cstring>
18 #include <filesystem>
19 #include <bits/stdc++.h>
20
21 const int MAX_LOG = 500;
22
23 // Hash function
24 struct vector_hash
25 {
26     size_t operator()(const std::vector<int>
27                       &myVector) const
28     {
29         std::hash<int> hasher;
30         size_t answer = 0;
31
32         for (int i : myVector)
33         {
34             answer ^= hasher(i) + 0x9e3779b9 +
35                       (answer << 6) + (answer >> 2);
36         }
37         return answer;
38     }
39 };
40
41 void get_msg_filename(int process_nr, char msg_filename[32]);
42
43 namespace Pet_kea
44 {
45     inline size_t SER_SIZE_CTRL_MSG_T(int recvd_cnt, int v_size)
46     {
47         return (6 * sizeof(int) + recvd_cnt * (sizeof(int) + v_size * sizeof(int)));
48     };
49     const int SER_VOID_SIZE = 2 * sizeof(int);
50
51     typedef enum message_type
52     {
53         MSG,
54         CTRL,
55         VOID,
56         COMM1,
57         COMM2,
58         COMM3,
59         COMM4
60     } msg_type;
61
62     struct fail_log_t
63     {
64         int id;
65         int fail_nr;
66         int res_time;
67         fail_log_t() {}
68         fail_log_t(int id, int fail_nr, int res_time) : id(id),
69                                                         fail_nr(fail_nr),
70                                                         res_time(res_time) {}
71     };
72
73     struct comm_msg_t

```

```

74 {
75     message_type msg_type;
76     int sending_process_nr;
77
78     int time_v_j;
79     std::vector<int> time_v_min;
80     int committed_cnt;
81     std::unordered_set<std::vector<int>, vector_hash> committed_msgs;
82 };
83
84 struct ctrl_msg_t
85 {
86     message_type msg_type;
87     int sending_process_nr;
88     struct fail_log_t log_entry;
89     int recieved_cnt;
90     std::set<std::pair<int, std::vector<int>>> recieved_msgs;
91 };
92
93 struct msg_t
94 {
95     message_type msg_type;
96     int sending_process_nr;
97     std::vector<int> time_v;
98     std::vector<int> fail_v;
99     int msg_size;
100     char *msg_buf;
101     ~msg_t()
102     {
103         free(msg_buf);
104     }
105 };
106
107 struct msg_log_t
108 {
109     int msg_size;
110     bool recipient;
111     int process_id;
112     char *msg_buf;
113     std::vector<int> time_v_sender;
114     std::vector<int> time_v_reciever;
115     std::vector<int> fail_v_sender;
116     int next_checkpoint;
117     ~msg_log_t()
118     {
119         free(msg_buf);
120     }
121     msg_log_t &operator=(const msg_log_t &other)
122     {
123         if (this != &other)
124         {
125             msg_size = other.msg_size;
126             recipient = other.recipient;
127             process_id = other.process_id;
128             time_v_sender = other.time_v_sender;
129             time_v_reciever = other.time_v_reciever;
130             fail_v_sender = other.fail_v_sender;
131             next_checkpoint = other.next_checkpoint;
132             char *temp_buf = (char *)malloc(msg_size);
133             memcpy(temp_buf, other.msg_buf, msg_size);
134             free(msg_buf);
135             msg_buf = temp_buf;
136         }
137         return *this;
138     }
139 };
140
141 /**
142  * @brief Prints the contents of a message.
143  * This function prints the contents of a message struct, including message type,
144  * sending process number, time vector, failure vector, message size.
145  * @param msg Pointer to the message struct to be printed.
146  */
147 void print_msg(struct msg_t *msg);
148
149 /**
150  * @brief Prints the contents of a control message.
151  * This function prints the contents of a control message struct, including message type,
152  * sending process number, log entry, received count, and received messages.
153  * @param msg Pointer to the control message struct to be printed.
154  */
155 void print_ctrl_msg(struct ctrl_msg_t *msg);
156
157 /**
158  * @brief Sends data to another process without it being recorded in the state.
159  * @param input Pointer to the data to be sent.
160  * @param fildes Array containing file descriptors of the pipe.
161  * @param size Size of the data to be sent.
162  * @return Number of bytes sent on success, -1 on failure.
163  */
164 int send_void(char *input, int fildes[2], int size);
165
166 class State
167 {
168 private:
169     int id;
170     std::vector<int> time_v;
171     std::vector<int> time_v_min;
172     std::vector<int> fail_v;
173     int **fildes;
174     msg_log_t *msg_log;
175     int msg_cnt;
176     std::vector<bool> commit_v;

```



```

177 std::vector<bool> remove_v;
178 std::vector<int> checkpoints;
179 std::vector<std::vector<int>> ck_time_v;
180
181 std::vector<struct fail_log_t> fail_log;
182 std::unordered_set<std::vector<int>, vector_hash> arrived_msgs;
183 std::unordered_set<std::vector<int>, vector_hash> arrived_ctrl;
184 std::unordered_set<std::vector<int>, vector_hash> committed_msg_set;
185 std::unordered_set<std::vector<int>, vector_hash> committed_recieve_events;
186 std::ofstream msg_out;
187 std::ofstream commit_out;
188
189 /**
190  * @brief Checks if a message is a duplicate message.
191  * @param msg Pointer to a msg_t structure representing the message to be checked.
192  * @return true if the message is a duplicate, false otherwise.
193  */
194 bool check_duplicate(struct msg_t *msg);
195
196 /**
197  * @brief Checks for duplicate control messages in the fail log.
198  * @param log The fail log structure containing control messages to be checked.
199  * @return True if duplicate control messages are found, false otherwise.
200  */
201 bool check_duplicate_ctrl(struct fail_log_t log);
202
203 /**
204  * @brief Checks for duplicate commit messages in the message log.
205  * @param l_msg Pointer to the message log structure containing commit messages to be checked.
206  * @return True if duplicate commit messages are found, false otherwise.
207  */
208 bool check_duplicate_commit(struct msg_log_t *l_msg);
209
210 /**
211  * @brief Checks if a message is orphaned.
212  * @param msg Pointer to a msg_t structure representing the message to be checked.
213  * @return true if the message is orphaned, false otherwise.
214  */
215 bool check_orphaned(struct msg_t *msg);
216
217 /**
218  * @brief Finds the next_checkpoint variable after removing certain checkpoints.
219  * @param removed_checkpoints A vector containing the IDs of checkpoints that have been removed.
220  * @param curr_next_checkpoint The ID of the current next_checkpoint variable before removal.
221  * @return The ID of the next checkpoint after considering the removed checkpoints.
222  */
223 int next_checkpoint_after_rem(std::vector<int> removed_checkpoints, int curr_next_checkpoint);
224
225 /**
226  * @brief Removes log entries based on provided indices.
227  * @param to_remove Vector containing indices of log entries to be removed.
228  * @param final_index Index of the final log entry in the log structure.
229  * @return New final index
230  */
231 int rem_log_entries(std::vector<int> to_remove, int final_index);
232
233 /**
234  * @brief Removes checkpoints.
235  * @param to_remove A vector containing the IDs of checkpoints to be removed.
236  */
237 void rem_checkpoints(std::vector<int> to_remove);
238
239 /**
240  * @brief Retrieves the next checkpoint.
241  * @param ptr A pointer to the current checkpoint.
242  * @return A pointer to the next checkpoint.
243  */
244 int *next_checkpoint(int *ptr);
245
246 /**
247  * @brief Serializes a commit message.
248  * @param msg Pointer to the commit message structure to be serialized.
249  * @param data Pointer to the character array where the serialized data will be stored.
250  */
251 void serialize_commit(struct comm_msg_t *msg, char *data);
252
253 /**
254  * @brief Deserializes a commit message.
255  * @param data Pointer to the character array containing the serialized commit message.
256  * @param msg Pointer to the commit message structure where the deserialized data will be stored.
257  */
258 void deserialize_commit(char *data, struct comm_msg_t *msg);
259
260 /**
261  * @brief Serializes a control message structure into a character array.
262  * @param msg Pointer to the control message structure to be serialized.
263  * @param data Pointer to the character array where the serialized data will be stored.
264  */
265 void serialize_ctrl(struct ctrl_msg_t *msg, char *data);
266
267 /**
268  * @brief Deserializes a control message from a character array.
269  * @param data Pointer to the character array containing serialized data.
270  * @param msg Pointer to the control message structure where deserialized data will be stored.
271  */
272 void deserialize_ctrl(char *data, struct ctrl_msg_t *msg);
273
274 /**
275  * @brief Serializes a general message structure into a character array.
276  * @param msg Pointer to the message structure to be serialized.
277  * @param data Pointer to the character array where the serialized data will be stored.
278  */
279 void serialize(struct msg_t *msg, char *data);

```

```

280 /**
281  * @brief Deserializes a general message from a character array.
282  * @param data Pointer to the character array containing serialized data.
283  * @param msg Pointer to the message structure where deserialized data will be stored.
284  */
285 void deserialize(char *data, struct msg_t *msg);
286
287 /**
288  * @brief Serializes a log message structure into a character array.
289  * @param log Pointer to the log message structure to be serialized.
290  * @param data Pointer to the character array where the serialized data will be stored.
291  */
292 void serialize_log(struct msg_log_t *log, char *data);
293
294 /**
295  * @brief Deserializes a log message from a character array.
296  * @param data Pointer to the character array containing serialized data.
297  * @param log Pointer to the log message structure where deserialized data will be stored.
298  * @return returns the size of the deserialized log message
299  */
300 int deserialize_log(char *data, struct msg_log_t *log);
301
302 /**
303  * @brief Serializes the state data.
304  * @param data A pointer to the character array where the serialized state data will be stored.
305  */
306 void serialize_state(char *data);
307
308 /**
309  * @brief Deserializes the state data.
310  * @param data A pointer to the character array containing the serialized state data.
311  */
312 void deserialize_state(char *data);
313
314 /**
315  * @brief Stores a general message.
316  * @param msg Pointer to the message structure to be stored.
317  * @param recipient -1 if the stored msg is a receive event,
318  *                  any positive integer for the recipient process_id.
319  * @return 0 on success, -1 on failure.
320  */
321 int store_msg(struct msg_t *msg, int recipient);
322
323 /**
324  * @brief Commits a list of messages.
325  * @param committed_msgs A vector containing the IDs of messages to be committed.
326  * @return Returns 0 if the messages are successfully committed.
327  *         Returns a non-zero value if an error occurs.
328  */
329 int commit_msgs(std::vector<int> committed_msgs);
330
331 /**
332  * @brief performs a rollback
333  * @param msg Pointer to the control message structure that activated the rollback
334  * @return 0 on success, -1 on failure.
335  */
336 int rollback(struct ctrl_msg_t *msg);
337
338 /**
339  * @brief Sends a COMM2 message to the target with the specified ID.
340  * @param target_id The ID of the target to which the commit message will be sent.
341  * @return Returns 0 if the commit message is successfully sent.
342  *         Returns a non-zero value if an error occurs.
343  */
344 int send_commit(int target_id);
345
346 /**
347  * @brief Commits all messages that can be safely committed
348  * @param is_instigator A boolean value indicating whether the calling function
349  *                     is the initiator of the commit procedure (true) or not (false).
350  * @return Returns 0 if the action is successfully committed.
351  *         Returns a non-zero value if an error occurs.
352  */
353 int commit(bool is_instigator);
354
355 /**
356  * @brief Removes all messages and checkpoints that can safely be removed
357  * @return Returns 0 if the data is successfully removed.
358  *         Returns a non-zero value if an error occurs.
359  */
360 int remove_data();
361
362 public:
363 const int SER_LOG_SIZE = 4 * sizeof(int) + time_v.size() * 3 * sizeof(int);
364 const int SER_MSG_SIZE = 3 * sizeof(int) + time_v.size() * 2 * sizeof(int);
365 const int SER_COMM1_SIZE = 3 * sizeof(int);
366 const int SER_COMM2_SIZE = 3 * sizeof(int);
367 inline int SER_COMM3_SIZE(int committed_cnt)
368 {
369     return (committed_cnt * time_v.size() * 2) * sizeof(int) + (3 + time_v.size()) * sizeof(int);
370 };
371 inline int SER_COMM4_SIZE(int committed_cnt)
372 {
373     return (committed_cnt * time_v.size() * 2) * sizeof(int) + 3 * sizeof(int);
374 };
375 inline int SER_STATE_SIZE(int arr_ctrl_size, int comm_msg_size, int comm_recv_events_size)
376 {
377     return (5 + (time_v.size() * 4) + ((comm_msg_size + comm_recv_events_size) * time_v.size() * 2) +
378             (3 * arr_ctrl_size)) * sizeof(int);
379 };
380
381 /**

```

```

382 * @brief Initializes to 0, when 0 no automatic checkpointing will be executed,
383 * otherwise if (NUM_stored_events % SAVE_CNT == 0) is true a checkpoint is made
384 */
385 int SAVE_CNT;
386
387 // Copy assignment operator
388 State &operator=(const State &other)
389 {
390     if (this != &other)
391     {
392         id = other.id;
393         time_v = other.time_v;
394         fail_v = other.fail_v;
395
396         if (fildes)
397         {
398             for (int i = 0; i < (int)time_v.size(); i++)
399             {
400                 free(fildes[i]);
401             }
402             free(fildes);
403         }
404         fildes = (int **)malloc((int)time_v.size() * sizeof(int *));
405         for (int i = 0; i < (int)time_v.size(); i++)
406         {
407             fildes[i] = (int *)malloc(2 * sizeof(int));
408             fildes[i][0] = other.fildes[i][0];
409             fildes[i][1] = other.fildes[i][1];
410         }
411
412         msg_cnt = other.msg_cnt;
413         msg_log_t *temp_log = (msg_log_t *)calloc(MAX_LOG, sizeof(msg_log_t));
414
415         for (int i = 0; i < other.msg_cnt; i++)
416         {
417             temp_log[i] = other.msg_log[i];
418         }
419         for (int i = other.msg_cnt - 1; i >= 0; i--)
420         {
421             std::vector<int>().swap(msg_log[i].time_v_reciever);
422             std::vector<int>().swap(msg_log[i].time_v_sender);
423             std::vector<int>().swap(msg_log[i].fail_v_sender);
424
425             free(msg_log[i].msg_buf);
426         }
427         free(msg_log);
428         msg_log = temp_log;
429
430         checkpoints = other.checkpoints;
431
432         ck_time_v = other.ck_time_v;
433
434         fail_log = other.fail_log;
435
436         arrived_msgs = other.arrived_msgs;
437
438         arrived_ctrl = other.arrived_ctrl;
439
440         msg_out.close();
441         char filename[32];
442         get_msg_filename(id, filename);
443         msg_out.open(filename,
444                     std::ofstream::out | std::ofstream::binary | std::ofstream::ate | std::ofstream::in);
445     }
446     return *this;
447 }
448
449 /**
450 * @brief Constructor for State class.
451 * @param process_nr The process number.
452 * @param process_cnt The total number of processes.
453 * @param fd The file descriptors from all processes.
454 * @param restart Flag indicating whether the process is being restarted.
455 */
456 State(int process_nr, int process_cnt, int (*fd)[2], bool restart);
457
458 /**
459 * @brief Destructor for State class.
460 */
461 ~State();
462
463 /**
464 * @brief Retrieves a message buffer.
465 * This function retrieves the message buffer at the specified index 'i'.
466 * @param i Index of the message buffer to retrieve.
467 * @return Pointer to the message buffer.
468 */
469 char *get_msg(int i);
470
471 /**
472 * @brief Retrieves the message log.
473 * This function retrieves the message log, which is an array of message buffers.
474 * @return Pointer to an array of message buffers.
475 */
476 char **get_msg_log();
477
478 /**
479 * @brief Creates a checkpoint of the process state and the recieved messages.
480 * @return 0 on success, -1 on failure.
481 */
482 int checkpoint();
483

```

```

484  /**
485   * @brief Signals a commit procedure to take place. the process that calls this function is the
         instigator
486   * of the commit procedure.
487   * @return Returns 0 upon successful signaling of the commit.
488   * Returns a non-zero value if an error occurs during signaling.
489   */
490  int signal_commit();
491
492  /**
493   * @brief Sends a message to another process that will be recorded in the state.
494   *
495   * @param input Pointer to the message to be sent.
496   * @param fildes Array containing file descriptors of the pipe.
497   * @param size Size of the message to be sent.
498   * @return Number of bytes sent on success, -1 on failure.
499   */
500  int send_msg(char *input, int process_id, int size);
501
502  /**
503   * @brief Receives a message from another process
504   * @param fildes Array containing file descriptors of the pipe.
505   * @param output Pointer to the buffer where the received message will be stored.
506   * @param size Size of the buffer.
507   * @return 0 on success, 1 after receiving a VOID msg, 2 after receiving a CTRL msg,
508   * 3 after receiving a duplicate or orphaned message that was discarded, -1 on failure.
509   */
510  int recv_msg(int fildes[2], char *output, int size);
511
512  /**
513   * @brief Sends control message to other processes to indicate that a failure occurred.
514   * @param fildes Array of arrays containing file descriptors of the pipes.
515   */
516  int send_ctrl();
517
518  /**
519   * @brief updates file descriptors
520   * @param process_id The process id of the file descriptors to be changed.
521   * @param fd The new file descriptors.
522   * @return 0 on success, -1 on failure.
523   */
524  int update_fd(int process_id, int fd[2]);
525  };
526 }
527
528 #endif

```

References

- [1] C. van Ek, “Optimistic recovery protocol for concurrent failures,” Ph.D. dissertation, Universiteit van Amsterdam, 2023.
- [2] W. R. Stevens, *UNIX Network Programming, Volume 2: Interprocess Communications*. Prentice Hall, 1999.
- [3] Anonymous, “Interprocess communication in distributed systems,” *GeeksforGeeks*, 2019, accessed: 2024-05-16. [Online]. Available: <https://www.geeksforgeeks.org/interprocess-communication-in-distributed-systems/>
- [4] —, “Inter-process communication, technique t1559,” *MITRE ATT&CK*, 2024, accessed: 2024-05-16. [Online]. Available: <https://attack.mitre.org/techniques/T1559/>
- [5] P. A. Bernstein and E. Newcomer, *Principles of transaction processing*. Morgan Kaufmann, 2009.
- [6] H. Garcia-Molina and K. Salem, “Sagas,” *ACM Sigmod Record*, vol. 16, no. 3, pp. 249–259, 1987.
- [7] S. L. Peterson and P. Kearns, “Rollback based on vector time,” in *Proceedings of 1993 IEEE 12th Symposium on Reliable Distributed Systems*. IEEE, 1993, pp. 68–77.
- [8] R. Koo and S. Toueg, “Checkpointing and rollback-recovery for distributed systems,” *IEEE Transactions on software Engineering*, no. 1, pp. 23–31, 1987.
- [9] R. A. Rutte, “Implementation of the peterson-kearns rollback recovery protocol,” 2024. [Online]. Available: https://github.com/robbybert/Peterson_kearns_rollback_algorithm
- [10] G. LLC, “Protobuf,” 2024. [Online]. Available: <https://protobuf.dev/>